

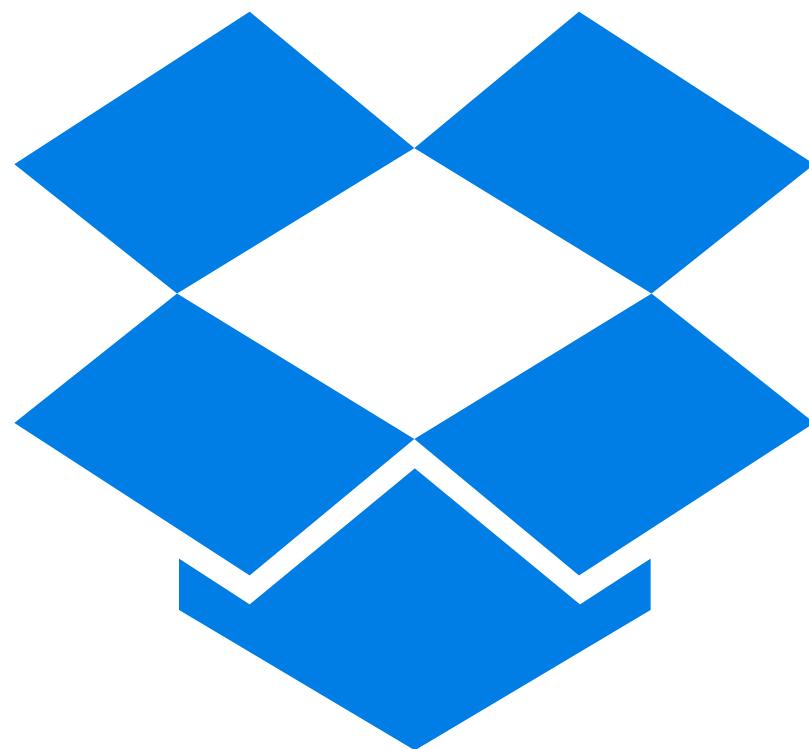
Bridging Languages Cross-Platform

Andrew Twyman & Jacob Potter

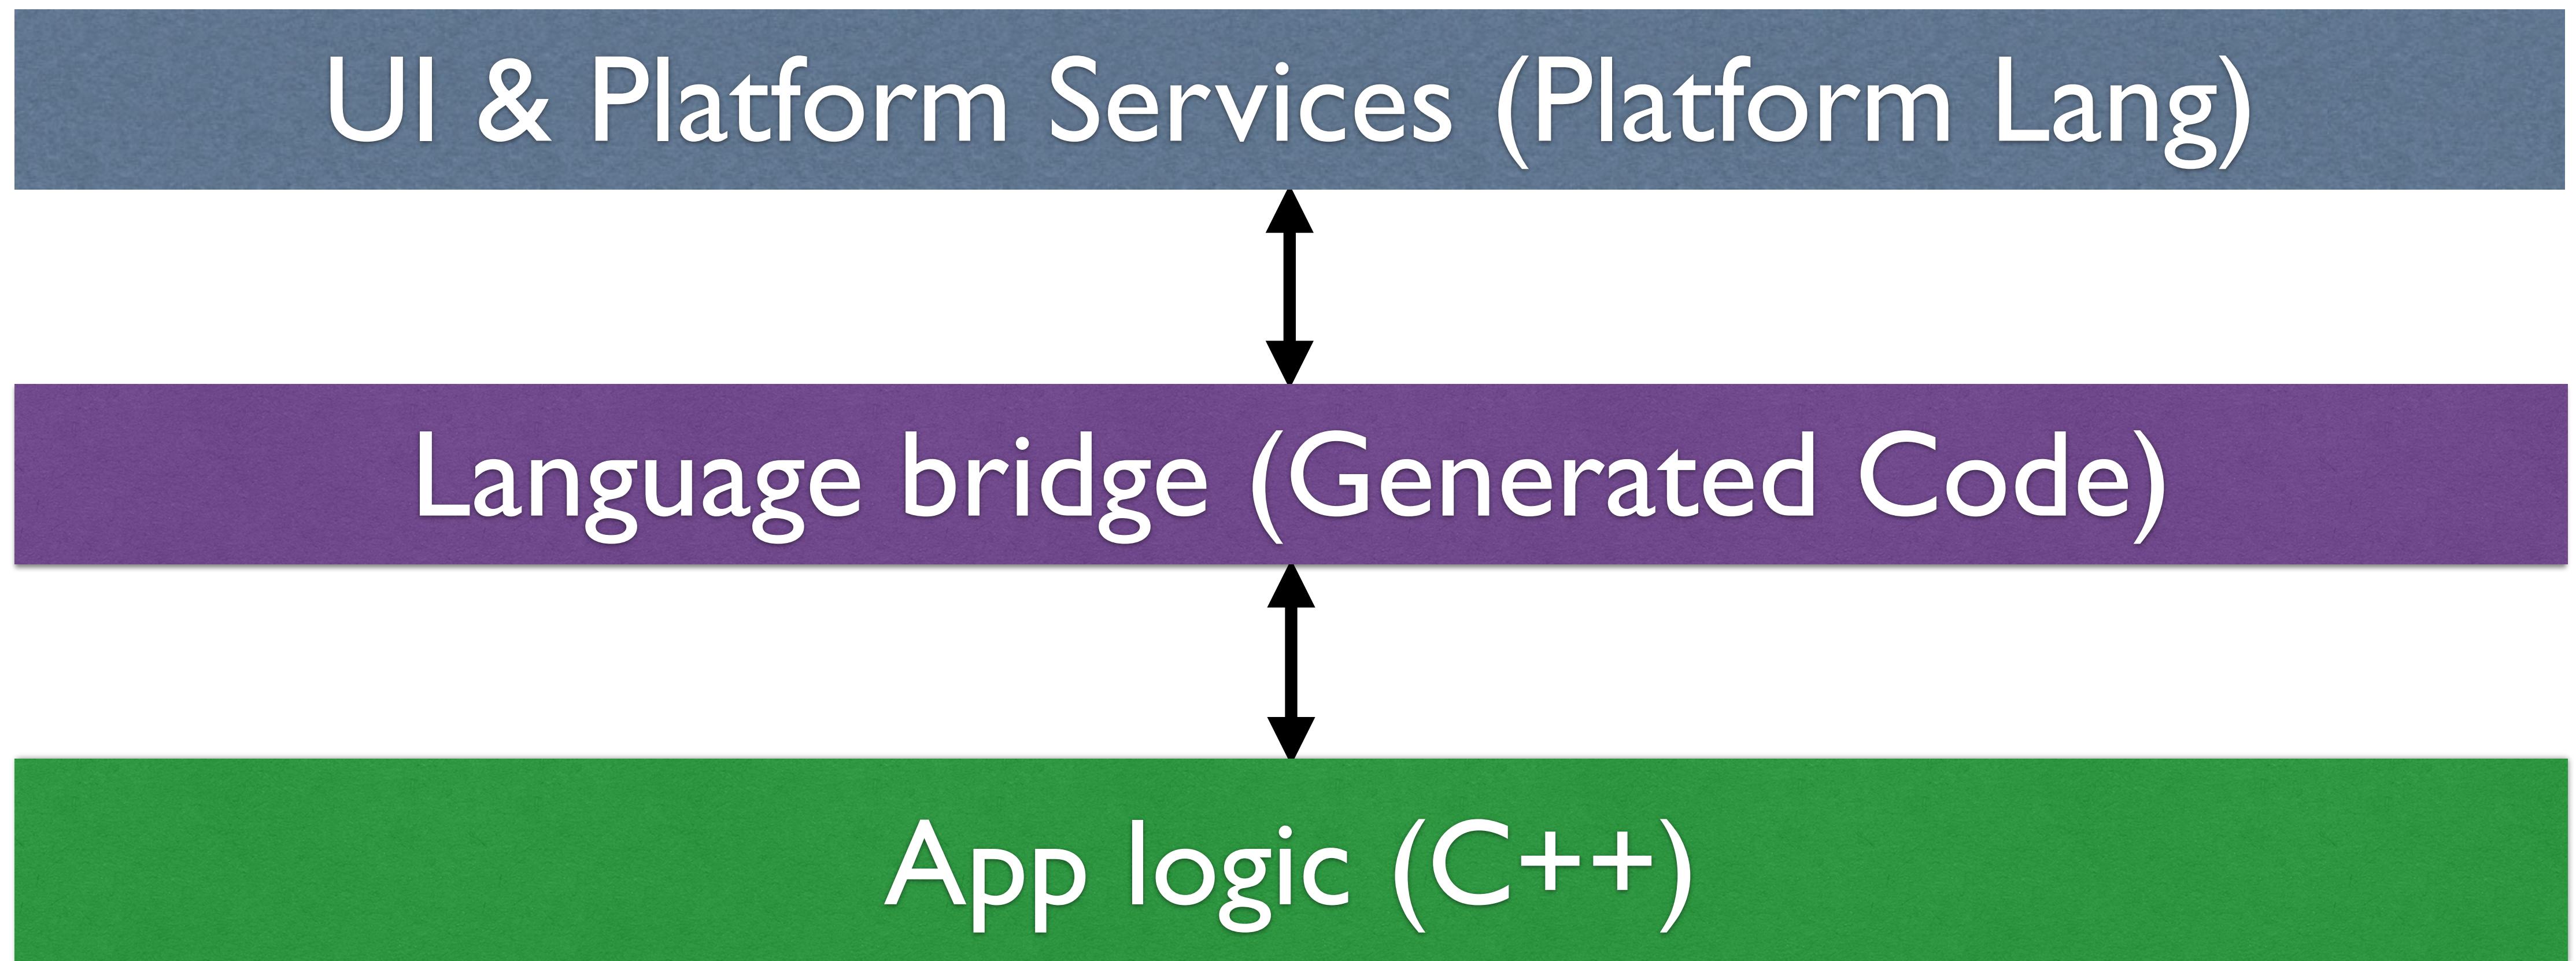


Who are we?

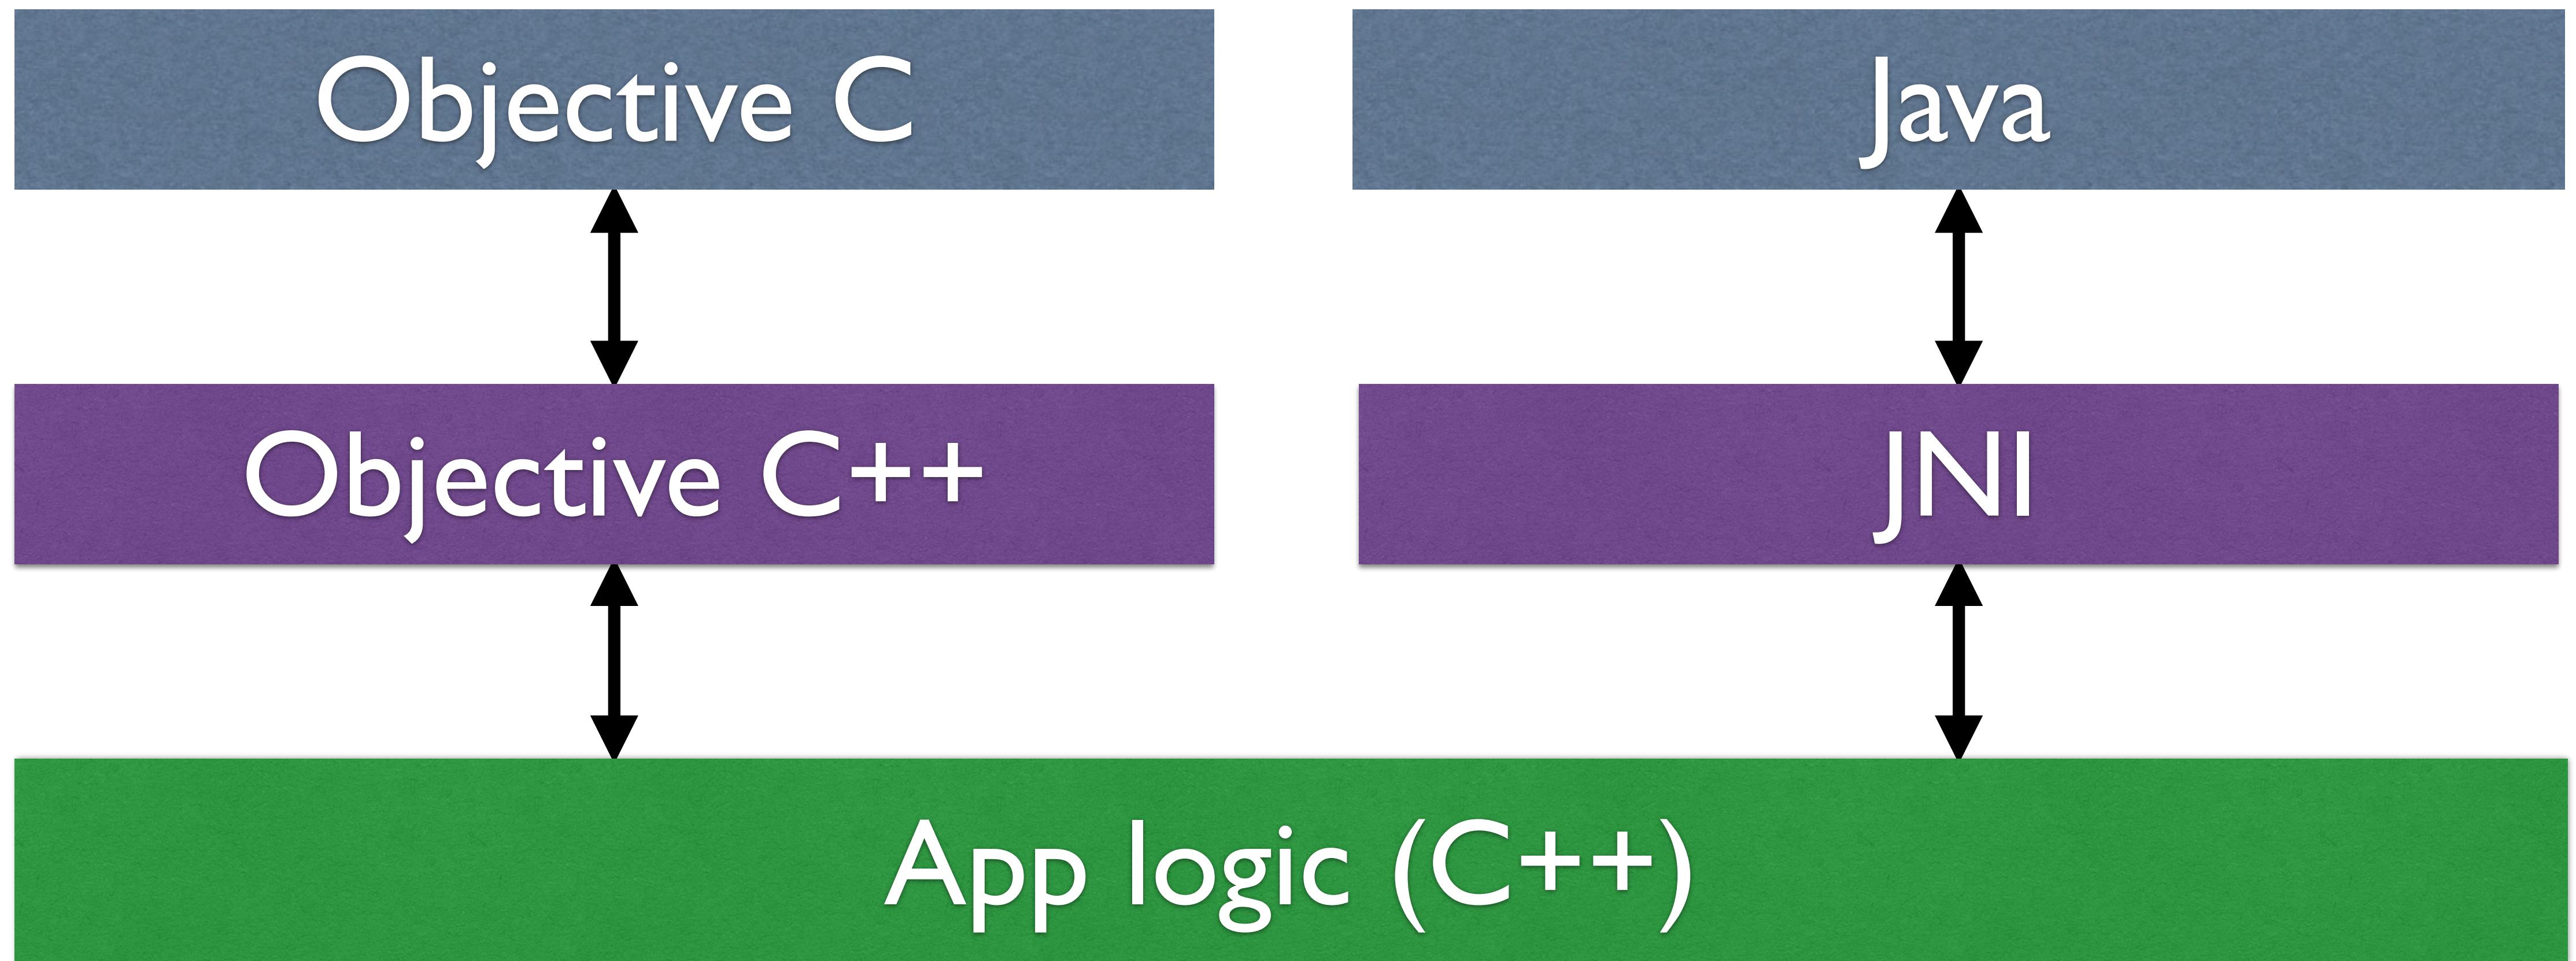
- ⌚ Cross-platform development at Dropbox
- ⌚ Mobile focused, but desktop and server are on our radar



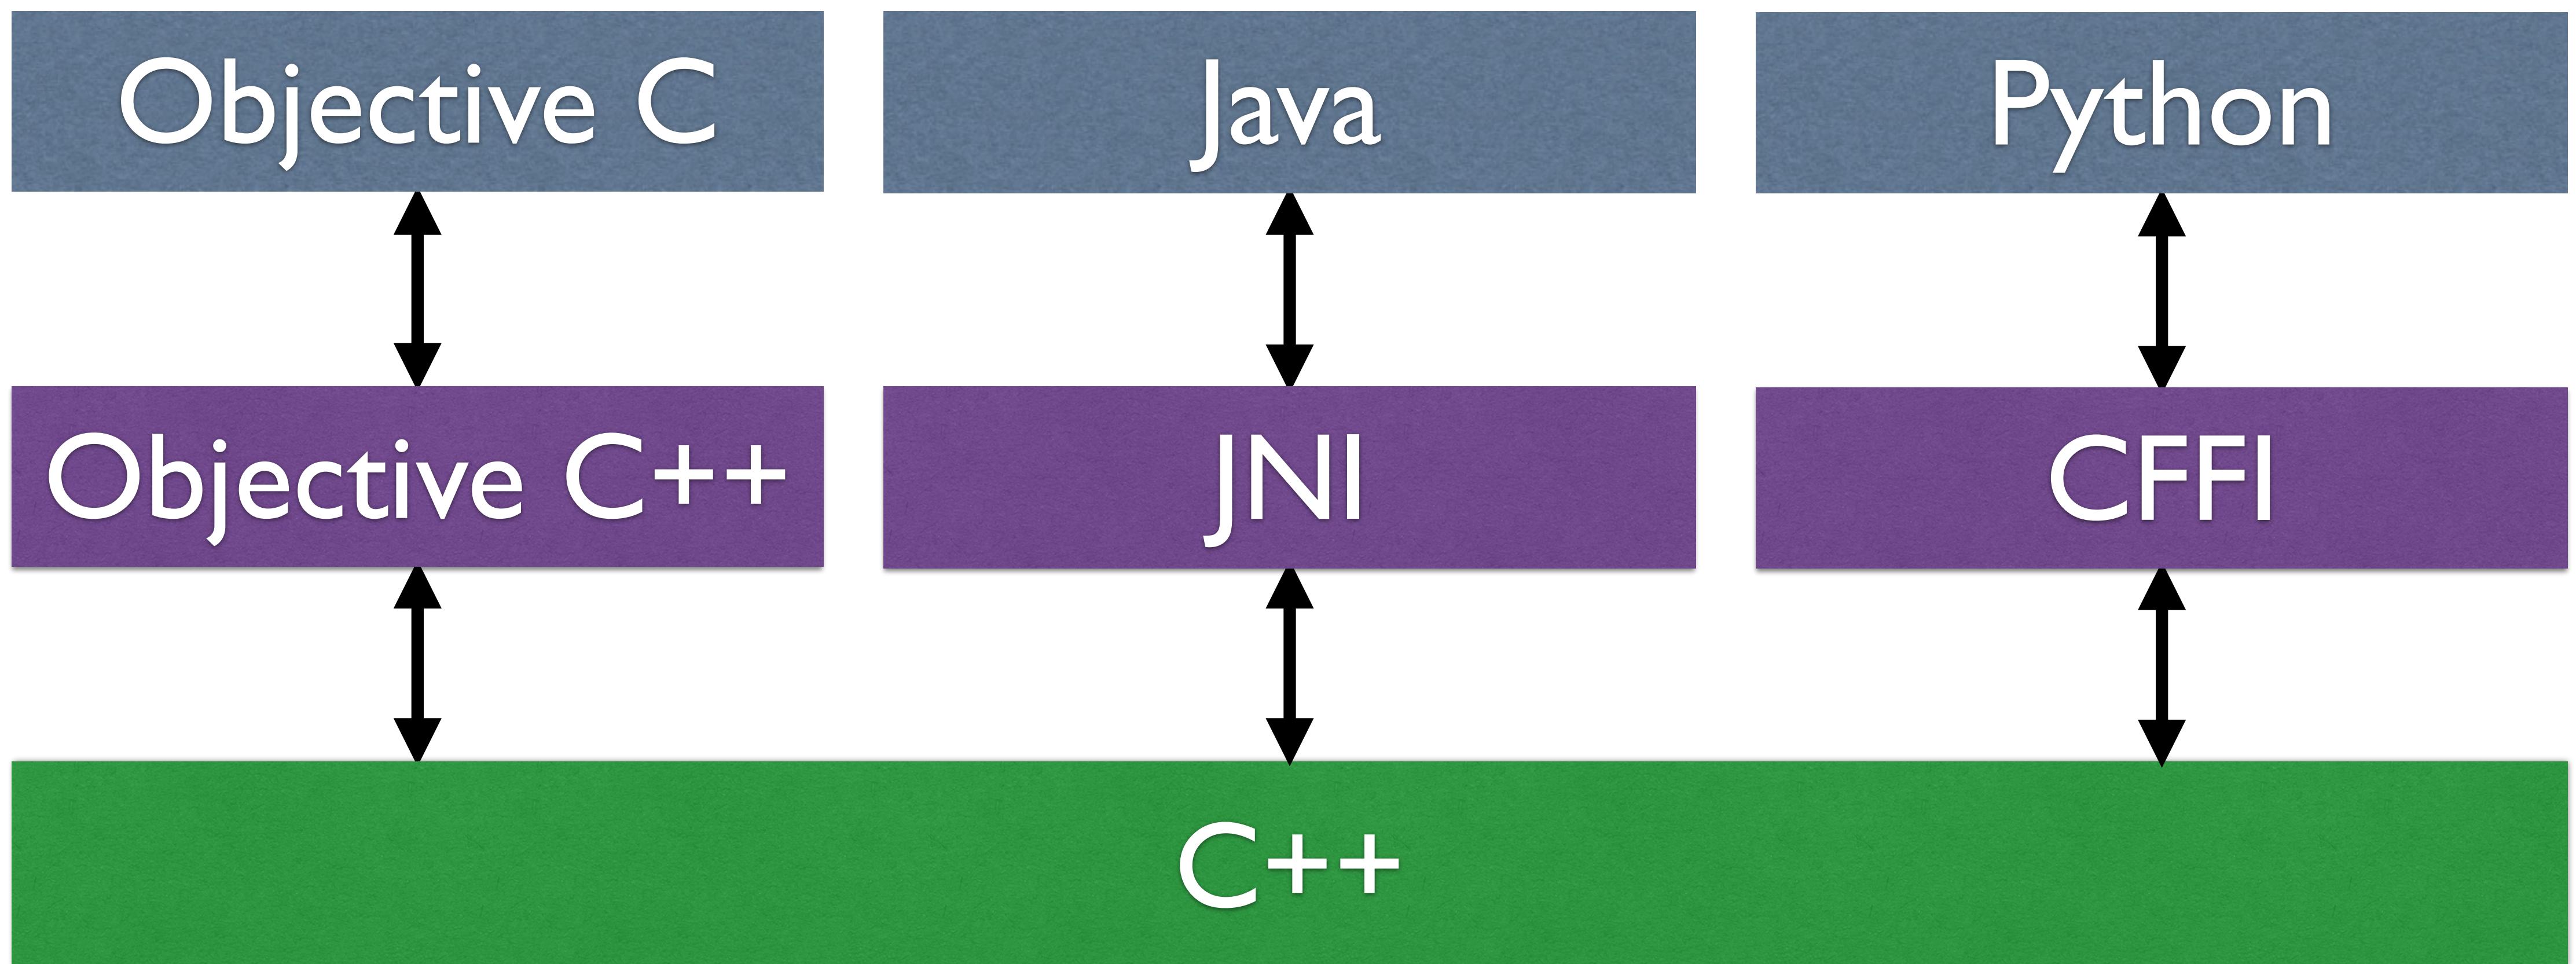
Our Cross-Platform Architecture



Supported Languages



Supported Languages



Today's non-Agenda

- ⌚ Not going to repeat last year's talk
 - ⌚ Why and how we created this architecture
 - ⌚ How we use Djinni to do it
 - ⌚ Check it out yourself:
 - ⌚ bit.ly/djinnivideo, bit.ly/djinnitalk
- ⌚ Not just about Djinni
 - ⌚ These are techniques you can use anywhere

Today's Agenda

- ⌚ Quick review: Language bridging with Djinni
- ⌚ Bridging a new language: Python
 - ⌚ I think our experience generalizes to other languages
- ⌚ Deep dives
 - ⌚ Interesting implementation techniques from all languages

What is Djinni for?

- ⌚ Say you're following our architecture, on Android & iOS
- ⌚ You want your UI code to call a common implementation
- ⌚ Your interface might be something like this

```
struct SomeInfo { /* ... */ } ;  
  
class MyModel {  
    void do_stuff(const SomeInfo & info);  
};
```

What do you do?

- ➊ (2 files) C++: Implement `MyModelImpl` class with one method
- ➋ (1 file) Java: Call from your UI code
- ➌ (1 file) ObjC: Call from your UI code
- ➍ (1 file) Djinni IDL: describes your interface

- ➎ With Djinni, you had to write 4 things
 - ➏ 3 of them are the custom code you have to write
 - ➐ 1 is IDL to define what you want the rest to be

What does Djinni do for you?

- ⌚ (3 files) MyModel abstract class (C++/Java/ObjC)
- ⌚ (3 files) SomeInfo data class (C++/Java/ObjC)
- ⌚ (4 files) JNI bridging: Java -> C++
- ⌚ (4 files) Objective-C++ bridging: ObjC -> C++

Djinni's Principles

- ⌚ Developers interact with natural-looking code in all languages
- ⌚ Method calls pass control between languages
 - ⌚ Not serialized RPC
- ⌚ Callable interfaces can be referenced across languages
- ⌚ Data is copied across languages

Djinni's Custom Types

- ➊ **Interfaces:** like abstract base classes
 - ➊ Implement in any language, own and call from any language
- ➋ **Records:** like structs
 - ➊ Immutable data, implemented for you, marshaled by copy
 - ➋ Contain data (not records) by value (no recursion)
- ➌ **Enums:** like scoped enums
 - ➊ Same values defined in all languages

Other stuff...

- ⌚ Primitive types: numbers, strings, bytes
- ⌚ Containers, Optionals
- ⌚ Static methods
- ⌚ Constants
- ⌚ Extensible and external types
- ⌚ ...
- ⌚ See GitHub for more

```
# An enum for specialized use.
wish_difficulty = enum {
    easy;
    medium;
    hard;
}

# A simple data struct.
wish = record {
    difficulty: wish_difficulty;
    request: string;
} deriving (eq, ord)

# A platform service, called by C++
djinni = interface +j +o {
    const max_wishes: i32 = 3;

    grant_wish(my_wish: wish): bool;
    past_wishes(): set<wish>;

    # Factory method
    static rub_lamp(): djinni;
}
```

Sync Progress

- ✓ Djinni overview
- ⌚ **Adding a new language (Python)**
 - ⌚ Problems presented by bridging to a new language
 - ⌚ Solutions we used for C++ to Python
- ⌚ Implementation techniques

Basic Steps of our Approach

- ➊ Pick a representation for IDL types
- ➋ Pick a bridging technology
- ➌ Does it meet the basic requirements?
- ➍ Build features on top of those basics

Step 1: Python Types

- ➊ Be idiomatic. Generated classes should look like Python.
- ➋ Use native Python types the programmer expects.
- ➌ Support the expected flexibility (duck typing).
- ➍ Support both Python 2, and Python 3.

Python Types: Basics

- ➊ Numbers: integer/long, float
 - ➊ Looser size constraints than Djinni
- ➋ Containers: list, dict, set
- ➌ Strings: unicode (2.7) or str (3.x)
- ➍ Bytes: str (2.7) or bytes (3.x)
- ➎ Optional: determines if None is legal

Python Types: Records

- ➊ Python class with named fields
- ➋ Python only has inclusion by reference
- ➌ `__init__` is all that defines the fields

```
class SomeInfo:  
    """ Info used by my data model. """  
  
    def __init__(self, user_id, obj_id, color):  
        self.user_id = user_id  
        self.obj_id = obj_id  
        self.color = color
```

Python Types: Enums

- ➊ Enum type introduced in Python 3.4
 - ➋ Available back-ported to Python 2
- ➌ IntEnum also acts like an int, for compatibility

```
@unique  
class Color(IntEnum):  
    Red = 0  
    Green = 1  
    Blue = 2
```

Python Types: Interfaces

- ⌚ Python class expected to have specific methods
- ⌚ Enforceable abstract base class introduced in Python 3
 - ⌚ Available back-ported to Python 2

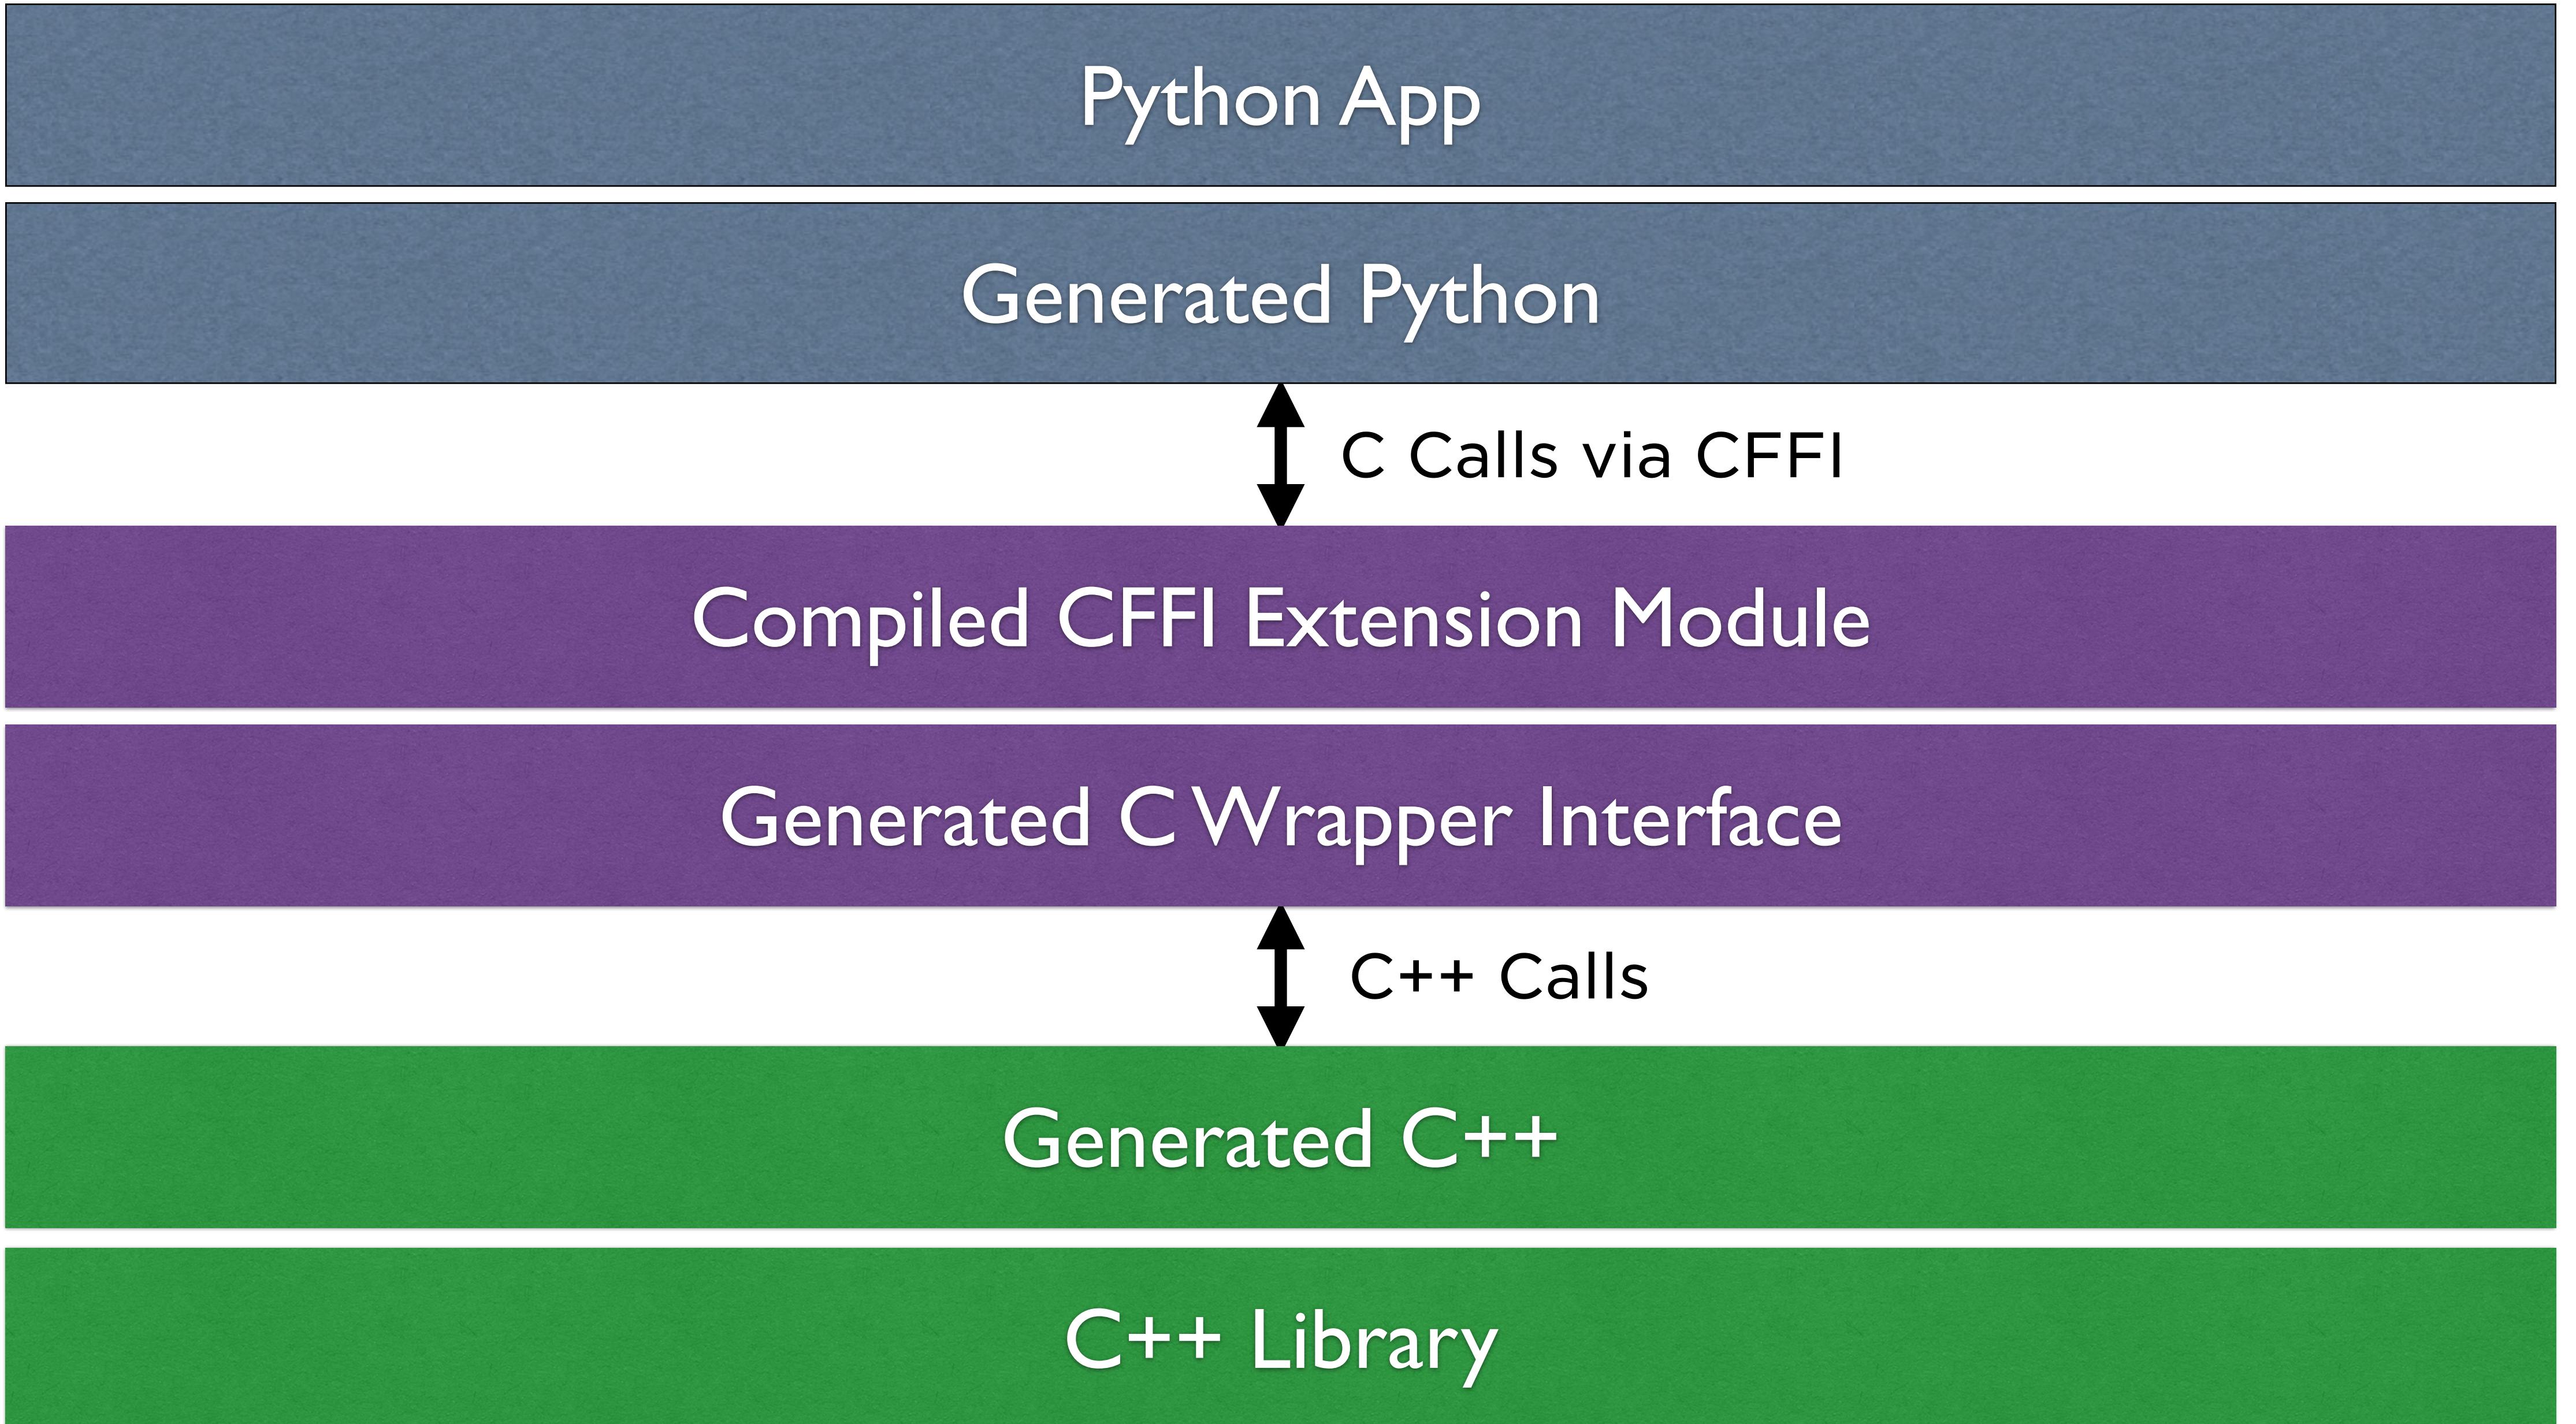
```
class MyModel(with_metaclass(ABCMeta)):  
    """ Represents the data model of my app. """  
  
    @abstractmethod  
    def do_stuff(self, info):  
        """ Does something with the given info. """  
        raise NotImplementedError
```

What about duck typing?

- ⌚ We C++ programmers like our strict type checks.
 - ⌚ Some Python programmers... not so much.
- ⌚ We provide tools so that user-defined code can **be** Djinni types.
- ⌚ But everything will work if user code **acts like** Djinni types.

Step 2: Pick a Bridging Tech

- ⌚ We use CFFI: “C Foreign Function Interface”
 - ⌚ Well accepted and supported
 - ⌚ Supports CPython & PyPy
 - ⌚ Supports Python 2 and 3
- ⌚ Generates efficient compiled code
- ⌚ Has multiple modes: we use “API Mode, Out-of-Line”



FFI at Build Time

- ➊ C++ code and C wrappers linked into a library (libfoo.dylib)
- ➋ Djinni feeds FFI declarations of our C wrapper functions

```
void cw__foo_setmsg(DjinniWrapperFoo * self,  
                     const char * msg)
```

- ➌ FFI generates C code to call them from Python
- ➍ Compiled into a Python extension (foo_cffi.so)
- ➎ Result: Compiled binary optimized for interacting with Python
 - ➏ Specific to a version and platform

CFFI at Run Time

- ➊ Load the compiled extension like any Python module

```
from foo_cffi import lib
```

- ➋ Call C functions from Python

```
lib.cw__foo_setmsg(cself, 'hello')
```

- ➌ Our generated Python classes hide this behind a proxy

```
foo.setmsg('hello')
```

Step 3: Basic Bridging

- ⌚ How do you call Python → C++?
- ⌚ How do you call C++ → Python?
- ⌚ How do you pass data?
- ⌚ How do you reference an object from the other language?

CFFI Calls: Python → C++

- ⌚ Easy after the build-time preparation

```
lib.cw__foo_setmsg(cself, 'hello')
```

FFI Calls: C++ → Python

- ⌚ Trickier since C++ can't speak to FFI directly
- ⌚ Python can create a callback for use by C++
- ⌚ Passed to C++ as function pointer (stored for later use)

```
@ffi.callback('void *(const char *)')
def log(msg):
    print(msg)

lib.cw__foo__addcallback_log(log)
```

- ⌚ We can do this during Python module load

CFFI Data

- ➊ CFFI can understand all C types
 - ➊ primitives, arrays, structs, pointers, ...
- ➋ We use:
 - ➊ primitives: bool, integers, floating-point
 - ➋ Pointers to opaque structs
- ➌ Wrapper structs are declared to CFFI but not defined
- ➍ All creation and manipulation of structs happens in C++

Objects: C++ → Python

- ➊ Wrap it in an opaque struct
- ➋ Pass a pointer to Python
- ➌ CFFI holds it as a `c_data` object

Objects: Python → C++

- ➊ `ffi.new_handle(obj)` creates a `c_data` from a Python object
- ➋ Pass to C++ as `void*` (actually we cast to `PyObjectHandle*`)
- ➌ C++ can pass it to a callback
- ➍ Recover the Python object with `ffi.from_handle(ptr)`
- ➎ Object lifetime is tricky - we'll come back to that

Step 4: Bridging Features

- ➊ Everything else can be built on those basics
- ➋ Primary required features for Djinni
 - ➌ Proxies for Interfaces
 - ➌ Marshaling structured data
 - ➌ Object ownership across languages
 - ➌ Handling exceptions

Proxy Objects

- ➊ A proxy is an interface object usable in language A, which stands in for an object in language B
 - ➊ Separate proxies for each direction
- ➋ Example: `MyModel` is an interface, so `MyModelCppProxy`
 - ➊ Is a Python object of Python class `MyModel`
 - ➊ Owns a reference to a C++ object of C++ class `MyModel`
 - ➊ Forwards all method calls across the bridge to the other object
- ➌ Deep dive into proxy management coming later

Marshaling Structured Data

- ➊ Marshaling takes multiple steps
 - ➊ Construct string
 - ➊ Construct record1
 - ➊ Put it into a list (and repeat)
 - ➊ Put that into record2
- ➋ But don't copy at every step

```
record1 = record {  
    s: string;  
}  
  
record2 = record {  
    list1: list<record1>;  
}
```

How do we minimize copies?

- ➊ C++ always controls the marshaling
 - ➊ C++ understands by-value and move
 - ➊ Python allows incremental construction of class fields
- ➋ Generate C++ to marshal a whole object either direction
- ⌁ Call into Python for sub-parts as necessary
- ⌂ Generator creates helper callbacks for each part

Python → C++

- ➊ Python passes reference to whole object
- ➋ C++ requests sub-parts as needed
 - ➌ One call per field of a record
 - ➌ Or iterator-style calls for a containers
- ➌ Constructs object all at once
- ➌ Move or RVO avoids copies between steps

C++ → Python

- ⌚ C++ asks Python to create an empty object
- ⌚ C++ passes sub-parts one at a time
 - ⌚ Add new class fields dynamically
 - ⌚ Or add to a list/dict/set
- ⌚ When done, pass full object to Python
- ⌚ Everything in Python is by-reference, so no copying

Object Ownership

- ⌚ No way to directly hold ownership between languages
- ⌚ C types cannot express ownership transfer
- ⌚ **Our golden rule:** An object passed across the language boundary comes with ownership
 - ⌚ (Except for a few special cases for internal helpers)

Implementing Ownership

- ➊ Once ownership is passed, it requires an explicit call to delete
- ➋ Automated via RAII wherever possible
 - ➌ unique_ptr with custom deleters
 - ➌ Helper classes and with blocks in Python

Explicit Ownership in Python

- ➊ Python's GC will try to delete an object we passed to C++
- ➋ Need to delay GC until C++ deletes it
- ➌ We use an extra global in Python
 - ➍ `ffi.new_handle(obj)` creates `c_data` object
 - ➎ Passed to C++ as a `void*`
 - ➏ Stored in a per-class `c_data_set` to keep it alive
 - ➐ Removed when a deleter callback is called

Explicit Ownership in Python

- ➊ handle = fi.new_handle(obj)
 - ➊ creates c_data object
- ➋ Passed to C++ as a void*
 - ➊ Python's GC doesn't track this.
 - ➊ Will delete the object! (Both obj and handle)
 - ➊ Need to delay GC until C++ deletes
- ➌ Store handle in a global per-class c_data_set to keep it alive
- ➍ Remove it in the deleter callback

Exclusive Ownership

- ➊ Most objects have unique ownership as they cross the bridge
 - ➊ Requires a callback to delete()
- ➋ C wrapper structs held by Python (all temporary)
 - ➊ DjinniString, DjinniBinary
 - ➋ BoxedI32, BoxedF64, etc.
- ➌ Python objects held by C++ (via handle)

Shared Ownership

- ➊ Interface objects need to cross the bridge more than once
 - ➊ Require ref-counting via calls to `inc_ref()` and `dec_ref()`
- ➋ C wrapper struct `DjinniWrapperMyModel`
 - ➊ Has its own ref-count for references from Python
 - ➊ Holds a `shared_ptr<MyModel>` for references from C++
 - ➊ `shared_ptr` might point to a C++ object, or a `PythonProxy`
 - ➊ Proxy holds unique ownership of a Python object (handle)

Propagating Exceptions

- ➊ Without help, CFFI will ignore exceptions, or crash
- ➋ Bridge code must catch, marshal, and re-throw
- ➌ C doesn't have exceptions, so we need a way to pass them
- ➍ We store the Exception in a per-thread state
 - ➎ Like errno or GetLastError()
 - ➏ But generated code will never forget to check

Exception State

- ➊ Per-thread variable holding an exception (or not)
 - ➋ Contains a Python exception
 - ➋ Held in C++ (by handle)
- ➋ Only populated while crossing the boundary
 - ➋ Checked and cleared after each call
- ➋ Symmetric code in both directions

Exception Responsibilities

- ➊ Callee
 - ➊ try/catch around full implementation
 - ➋ In catch block, set exception state and return
- ➋ Caller
 - ➊ After return, check exception state
 - ➋ If found, clear state and throw

Marshaling Exceptions

- ➊ Djinni provides simple translation by default
- ➋ You can plug in your own custom version
- ➌ The model is largely the same as Java
 - ➍ I'll describe it in a deep dive later

Python Wrap-up

- ➊ The rest can be built on this foundation
 - ➊ Hand-written marshaling helpers
 - ➋ Global proxy cache
 - ➋ Derived comparison/hash operations
 - ➋ Static methods, constants, ...
- ➋ Python is mostly complete, but still experimental
 - ➊ Check it out on the python branch on GitHub
 - ➋ <https://github.com/dropbox/djinni/tree/python>

Sync Progress

- ✓ Djinni overview
- ✓ Adding a new language (Python)
- ⌚ Implementation techniques
 - ⌚ **Proxy caching for identity semantics**
 - ⌚ Nullability
 - ⌚ Customized Java exception translation

Proxy objects

- ➊ We generate interfaces in each language
- ➋ Also generate implementations that perform cross-language calls
 - ➌ Java calling C++
 - ➍ `native void foo();`
 - ➎ `void Java_pkg_Widget_foo(JNIEnv *, jobject) { ... }`
 - ➏ C++ calling Java

```
weather = record {
    . . .
}

# An object that receives weather reports
weather_listener = interface +j +o {
    weather_report(date: date, forecast: weather);
}

# A service for subscribing to weather reports
weather_service = interface +c {
    add_listener(listener: weather_listener);
    remove_listener(listener: weather_listener);
    # ...
}
```

```
class MyWeatherService : public WeatherService
{
public:
    void add_listener(const shared_ptr<WeatherListener> & listener) {
        m_listeners.insert(listener);
    }
    void remove_listener(const shared_ptr<WeatherListener> & listener) {
        m_listeners.erase(listener);
    }

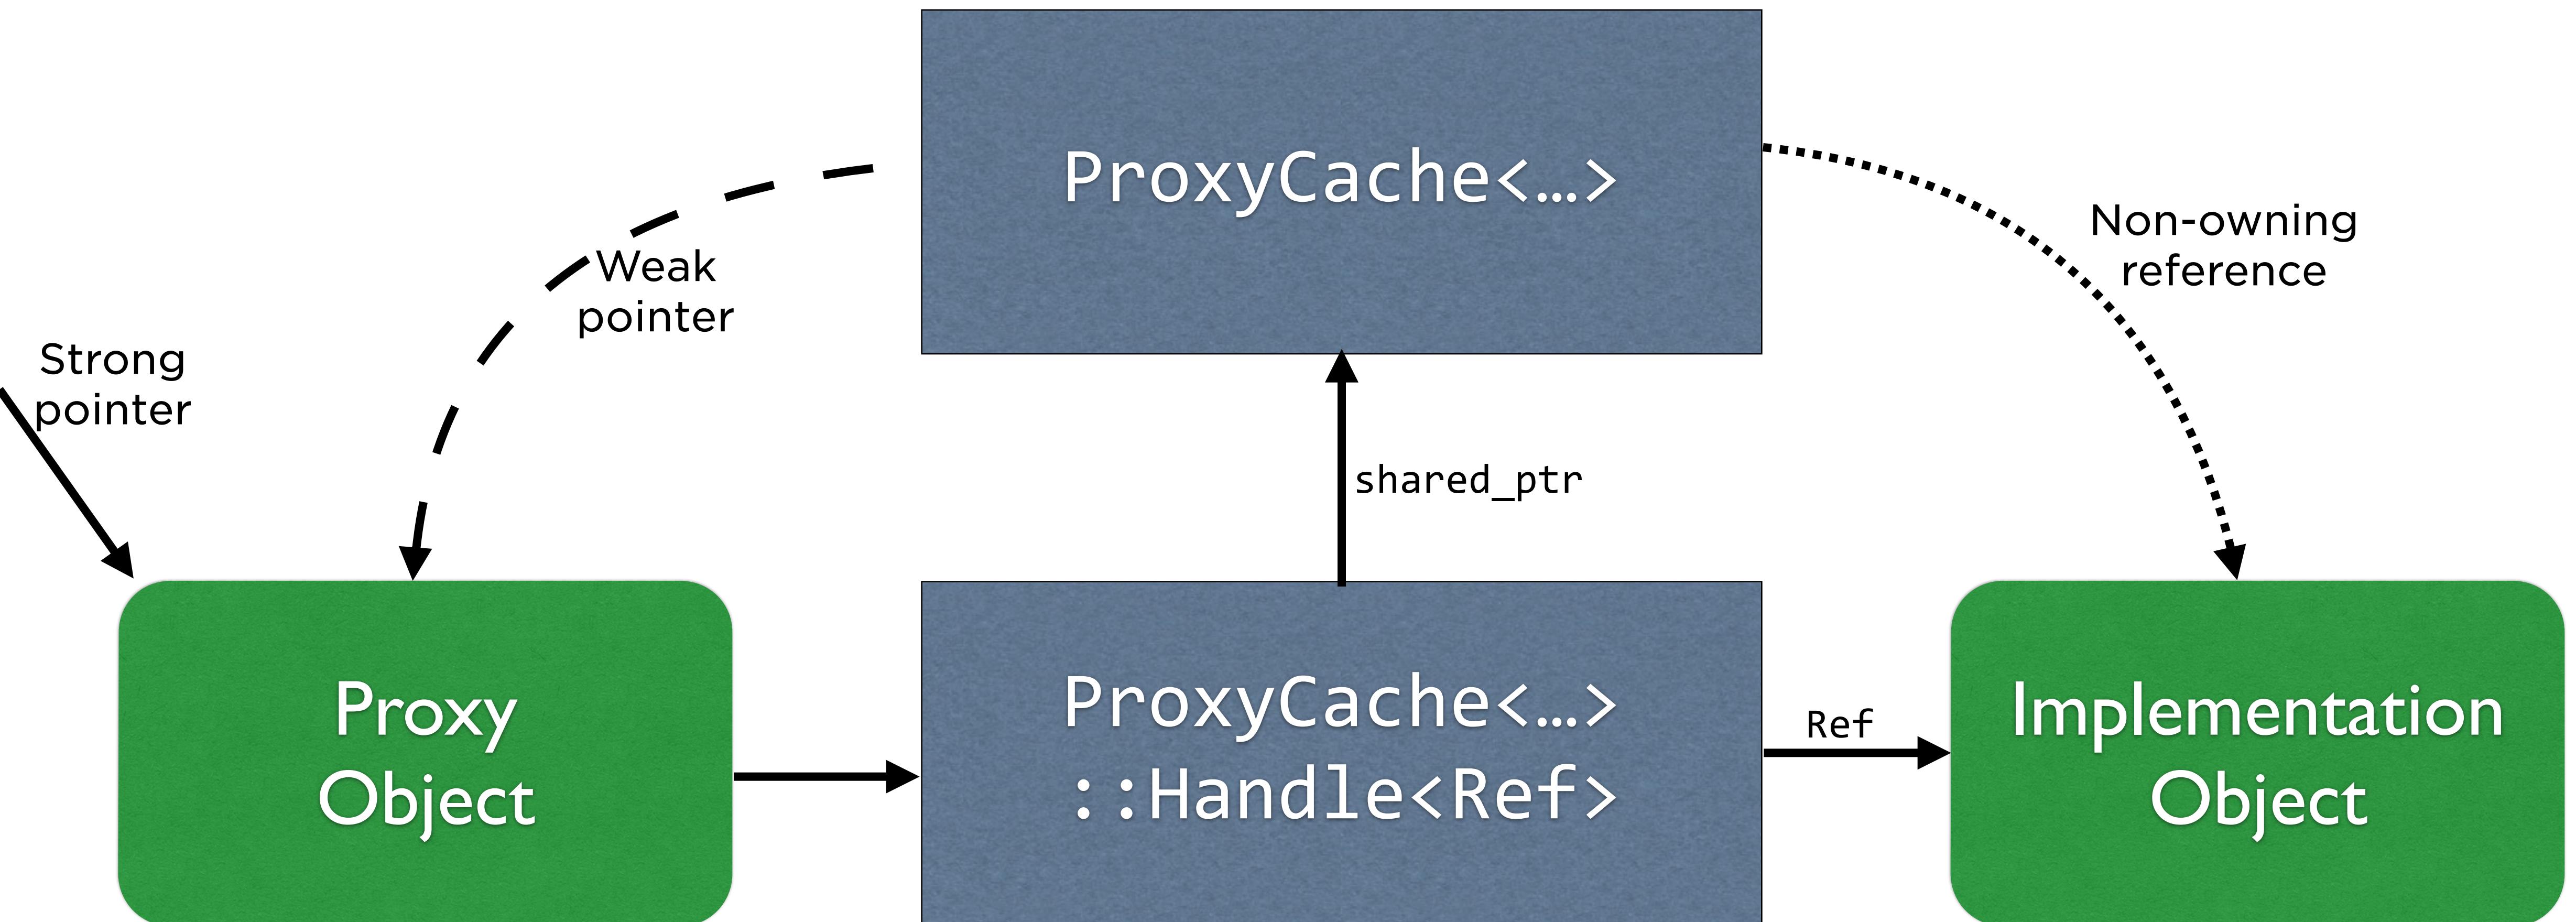
private:
    set<shared_ptr<WeatherListener>> m_listeners;
};
```

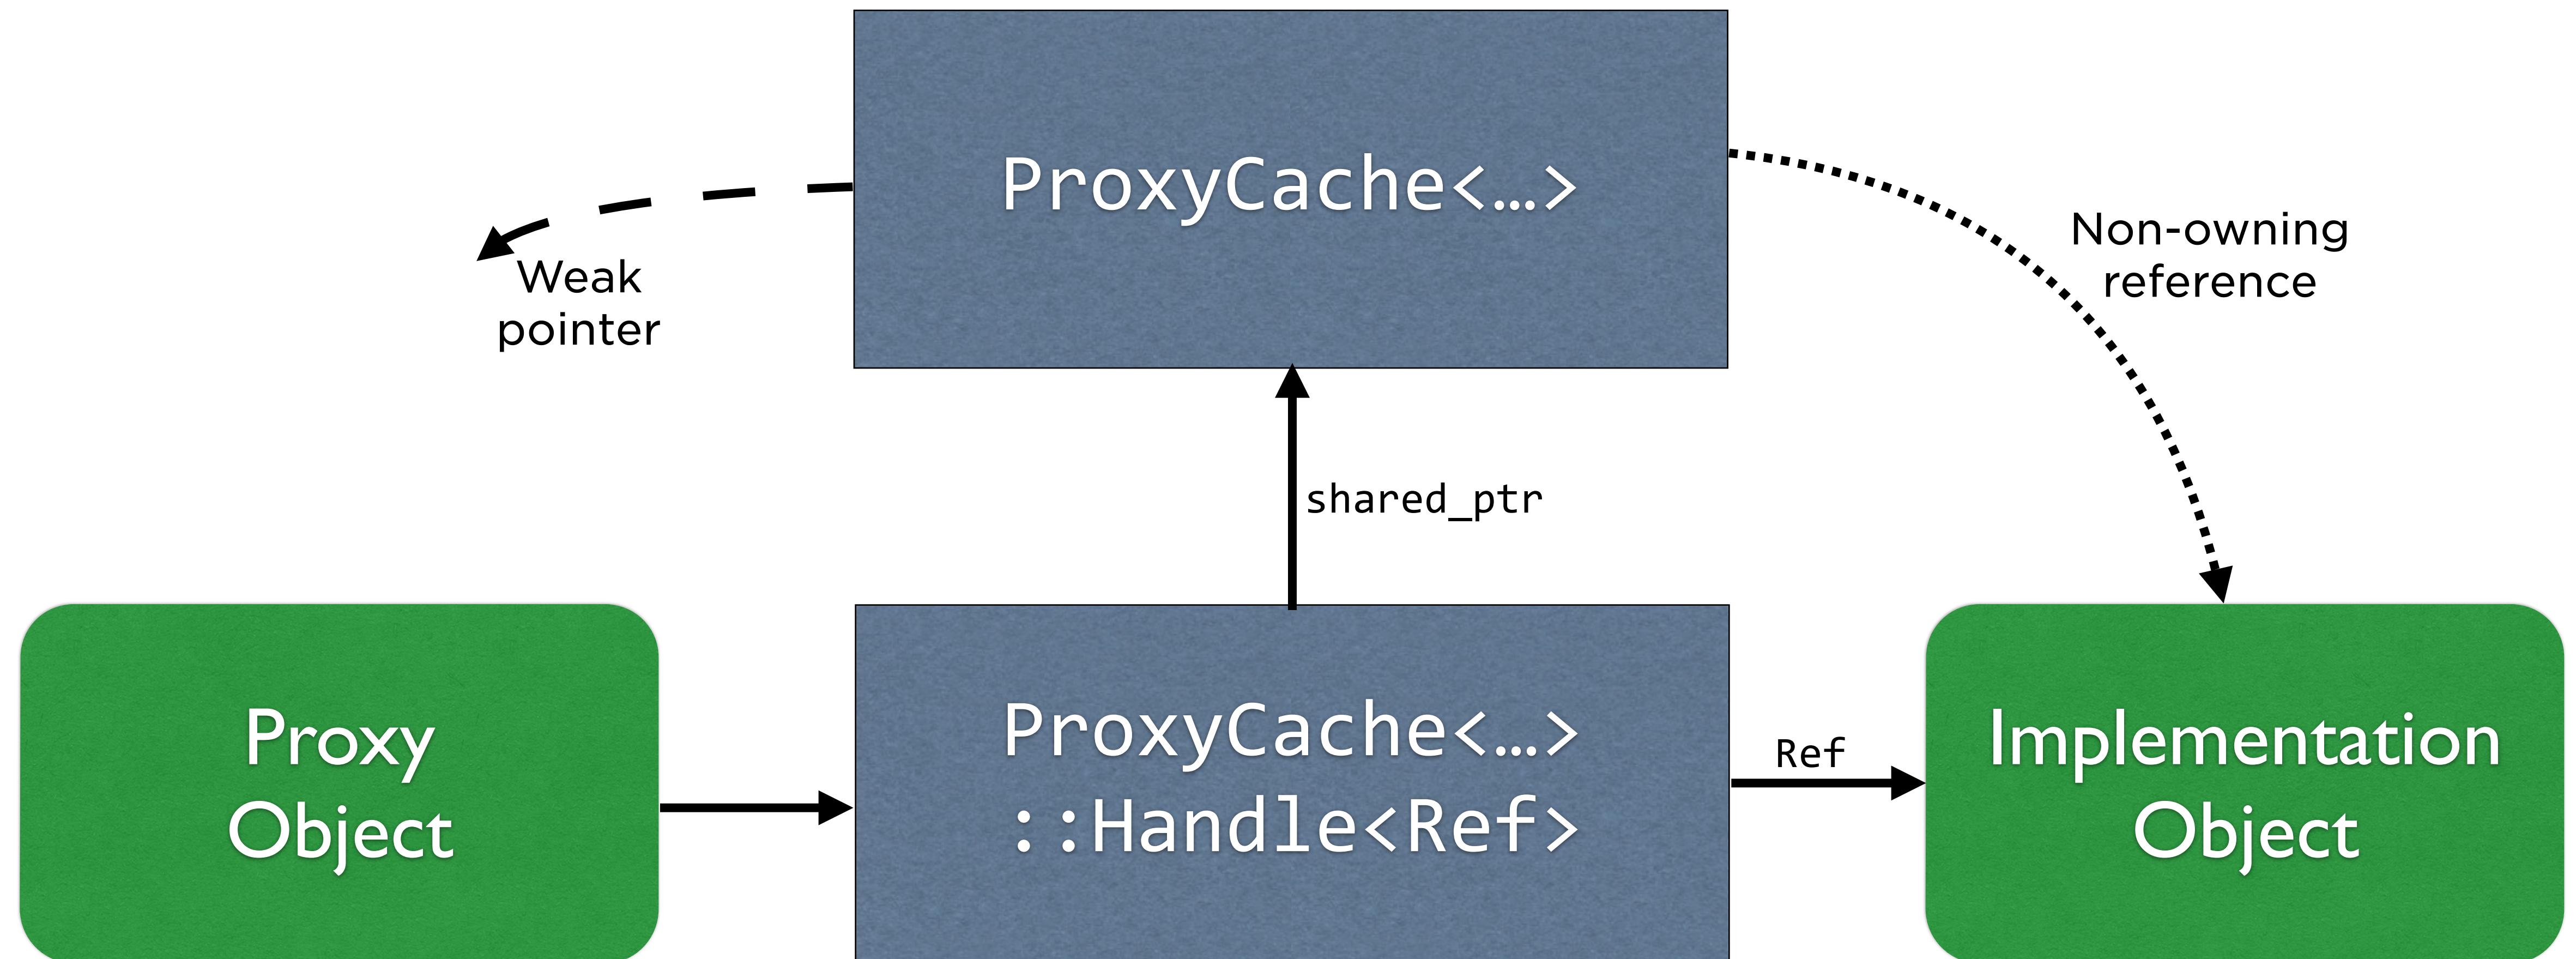
Proxy caching

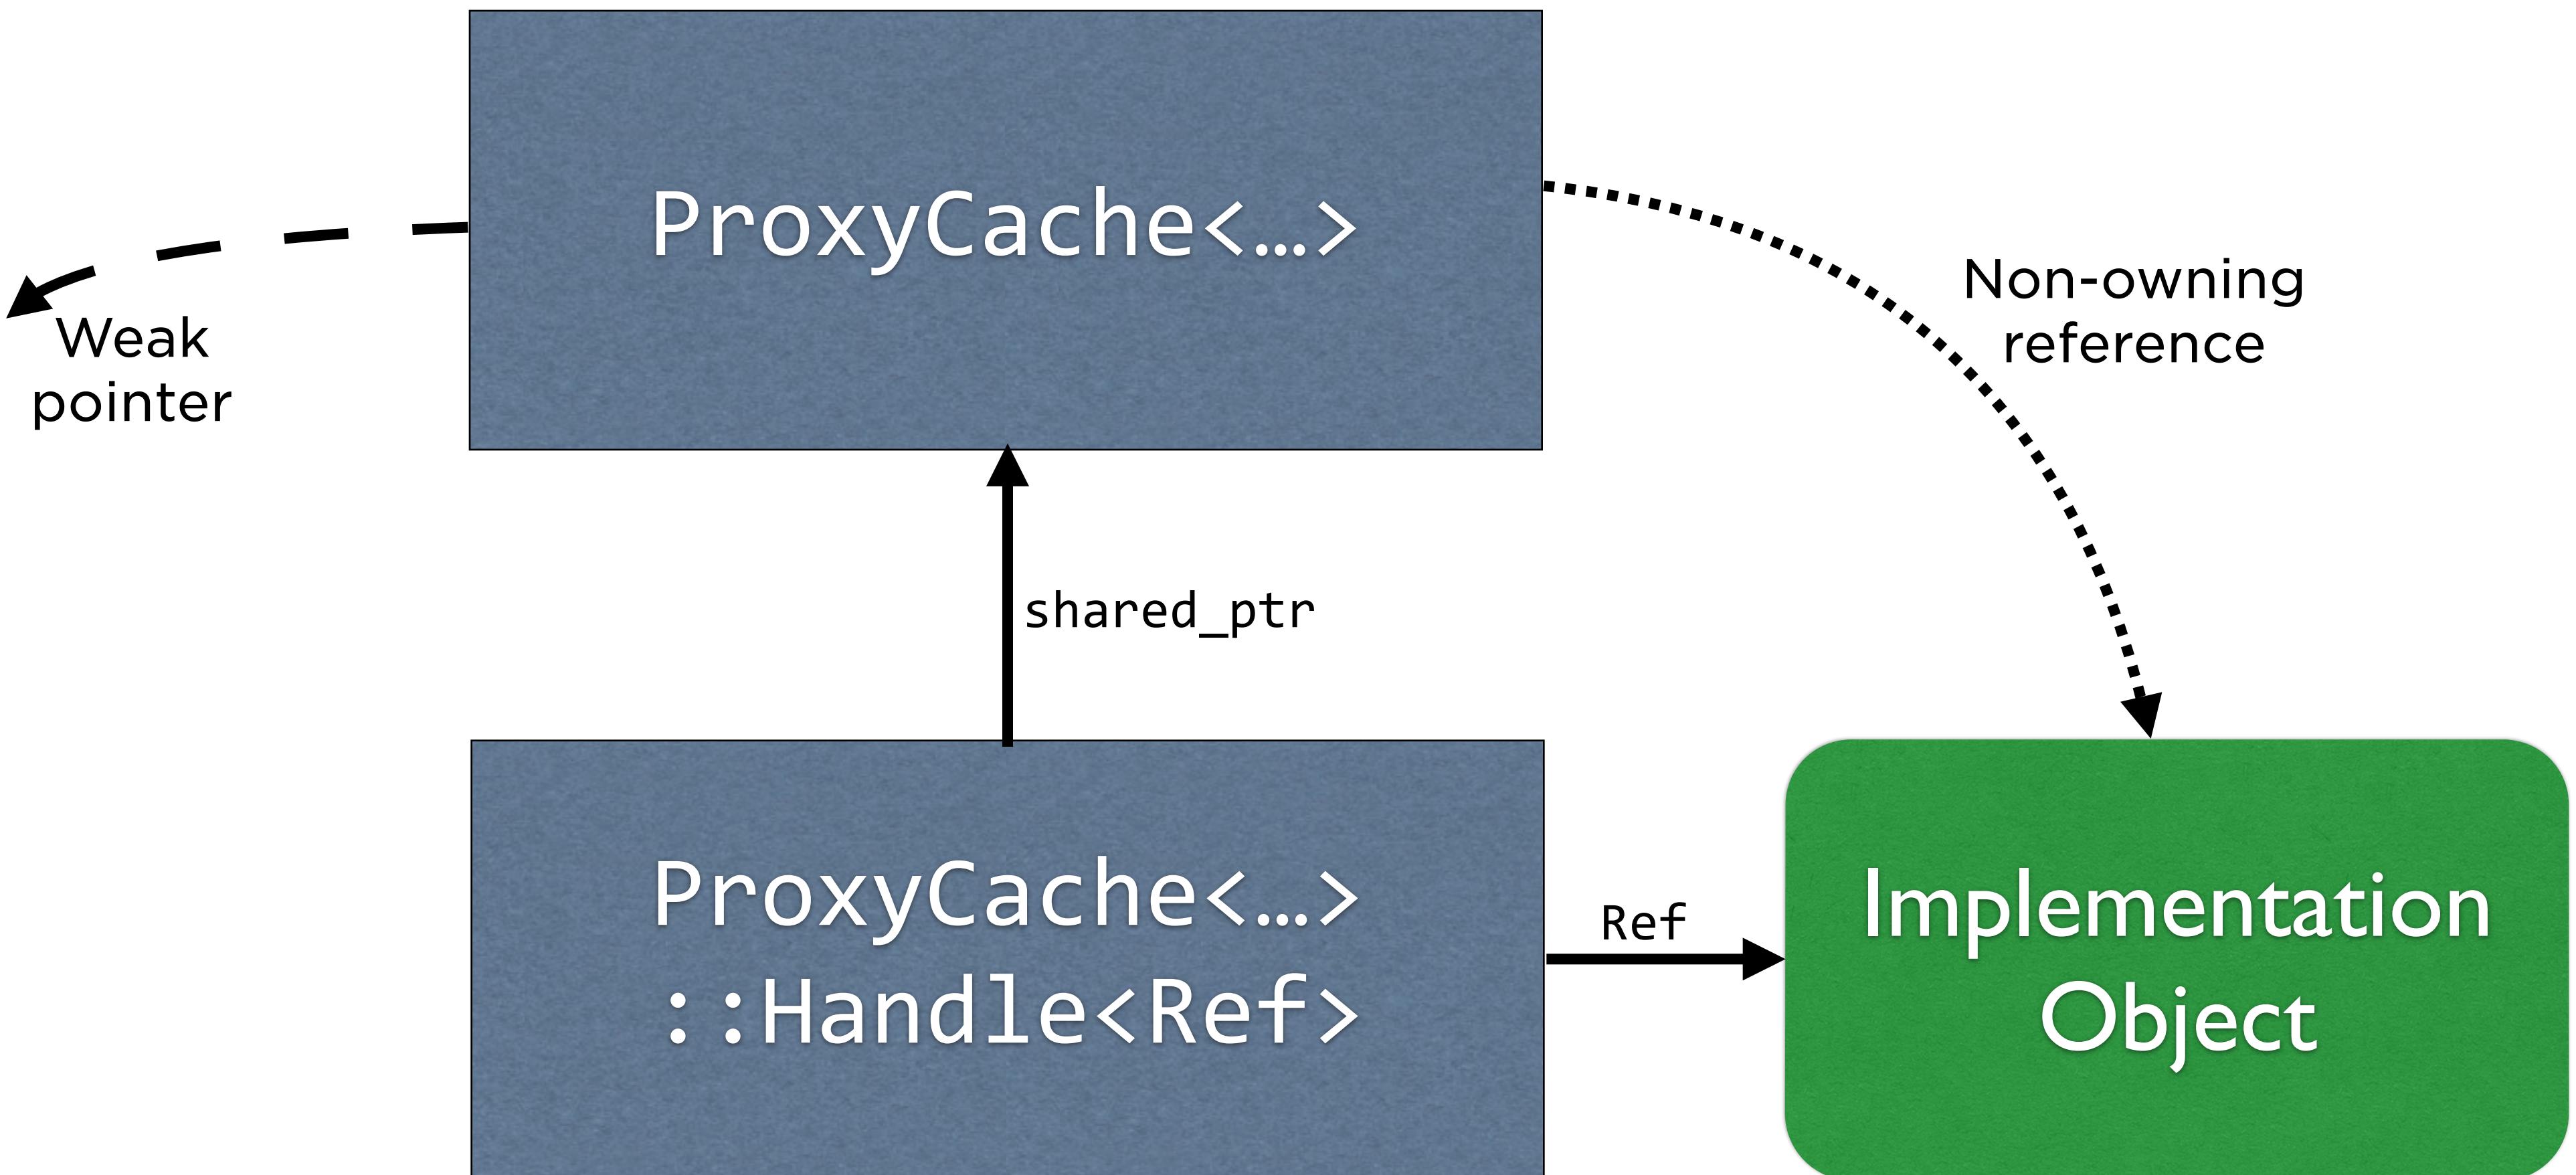
- ➊ When we call `set::erase`, we need the same listener object, even though we made two different calls through the bridge
- ➋ Performance benefits too, but mostly this is for correctness
- ➌ Need to enable idioms that developers want to use
- ➍ Solution:
 - ➎ Keep a reference to proxy objects, and reuse if possible
 - ➏ Store map from *non-owning reference* to *weak pointer*

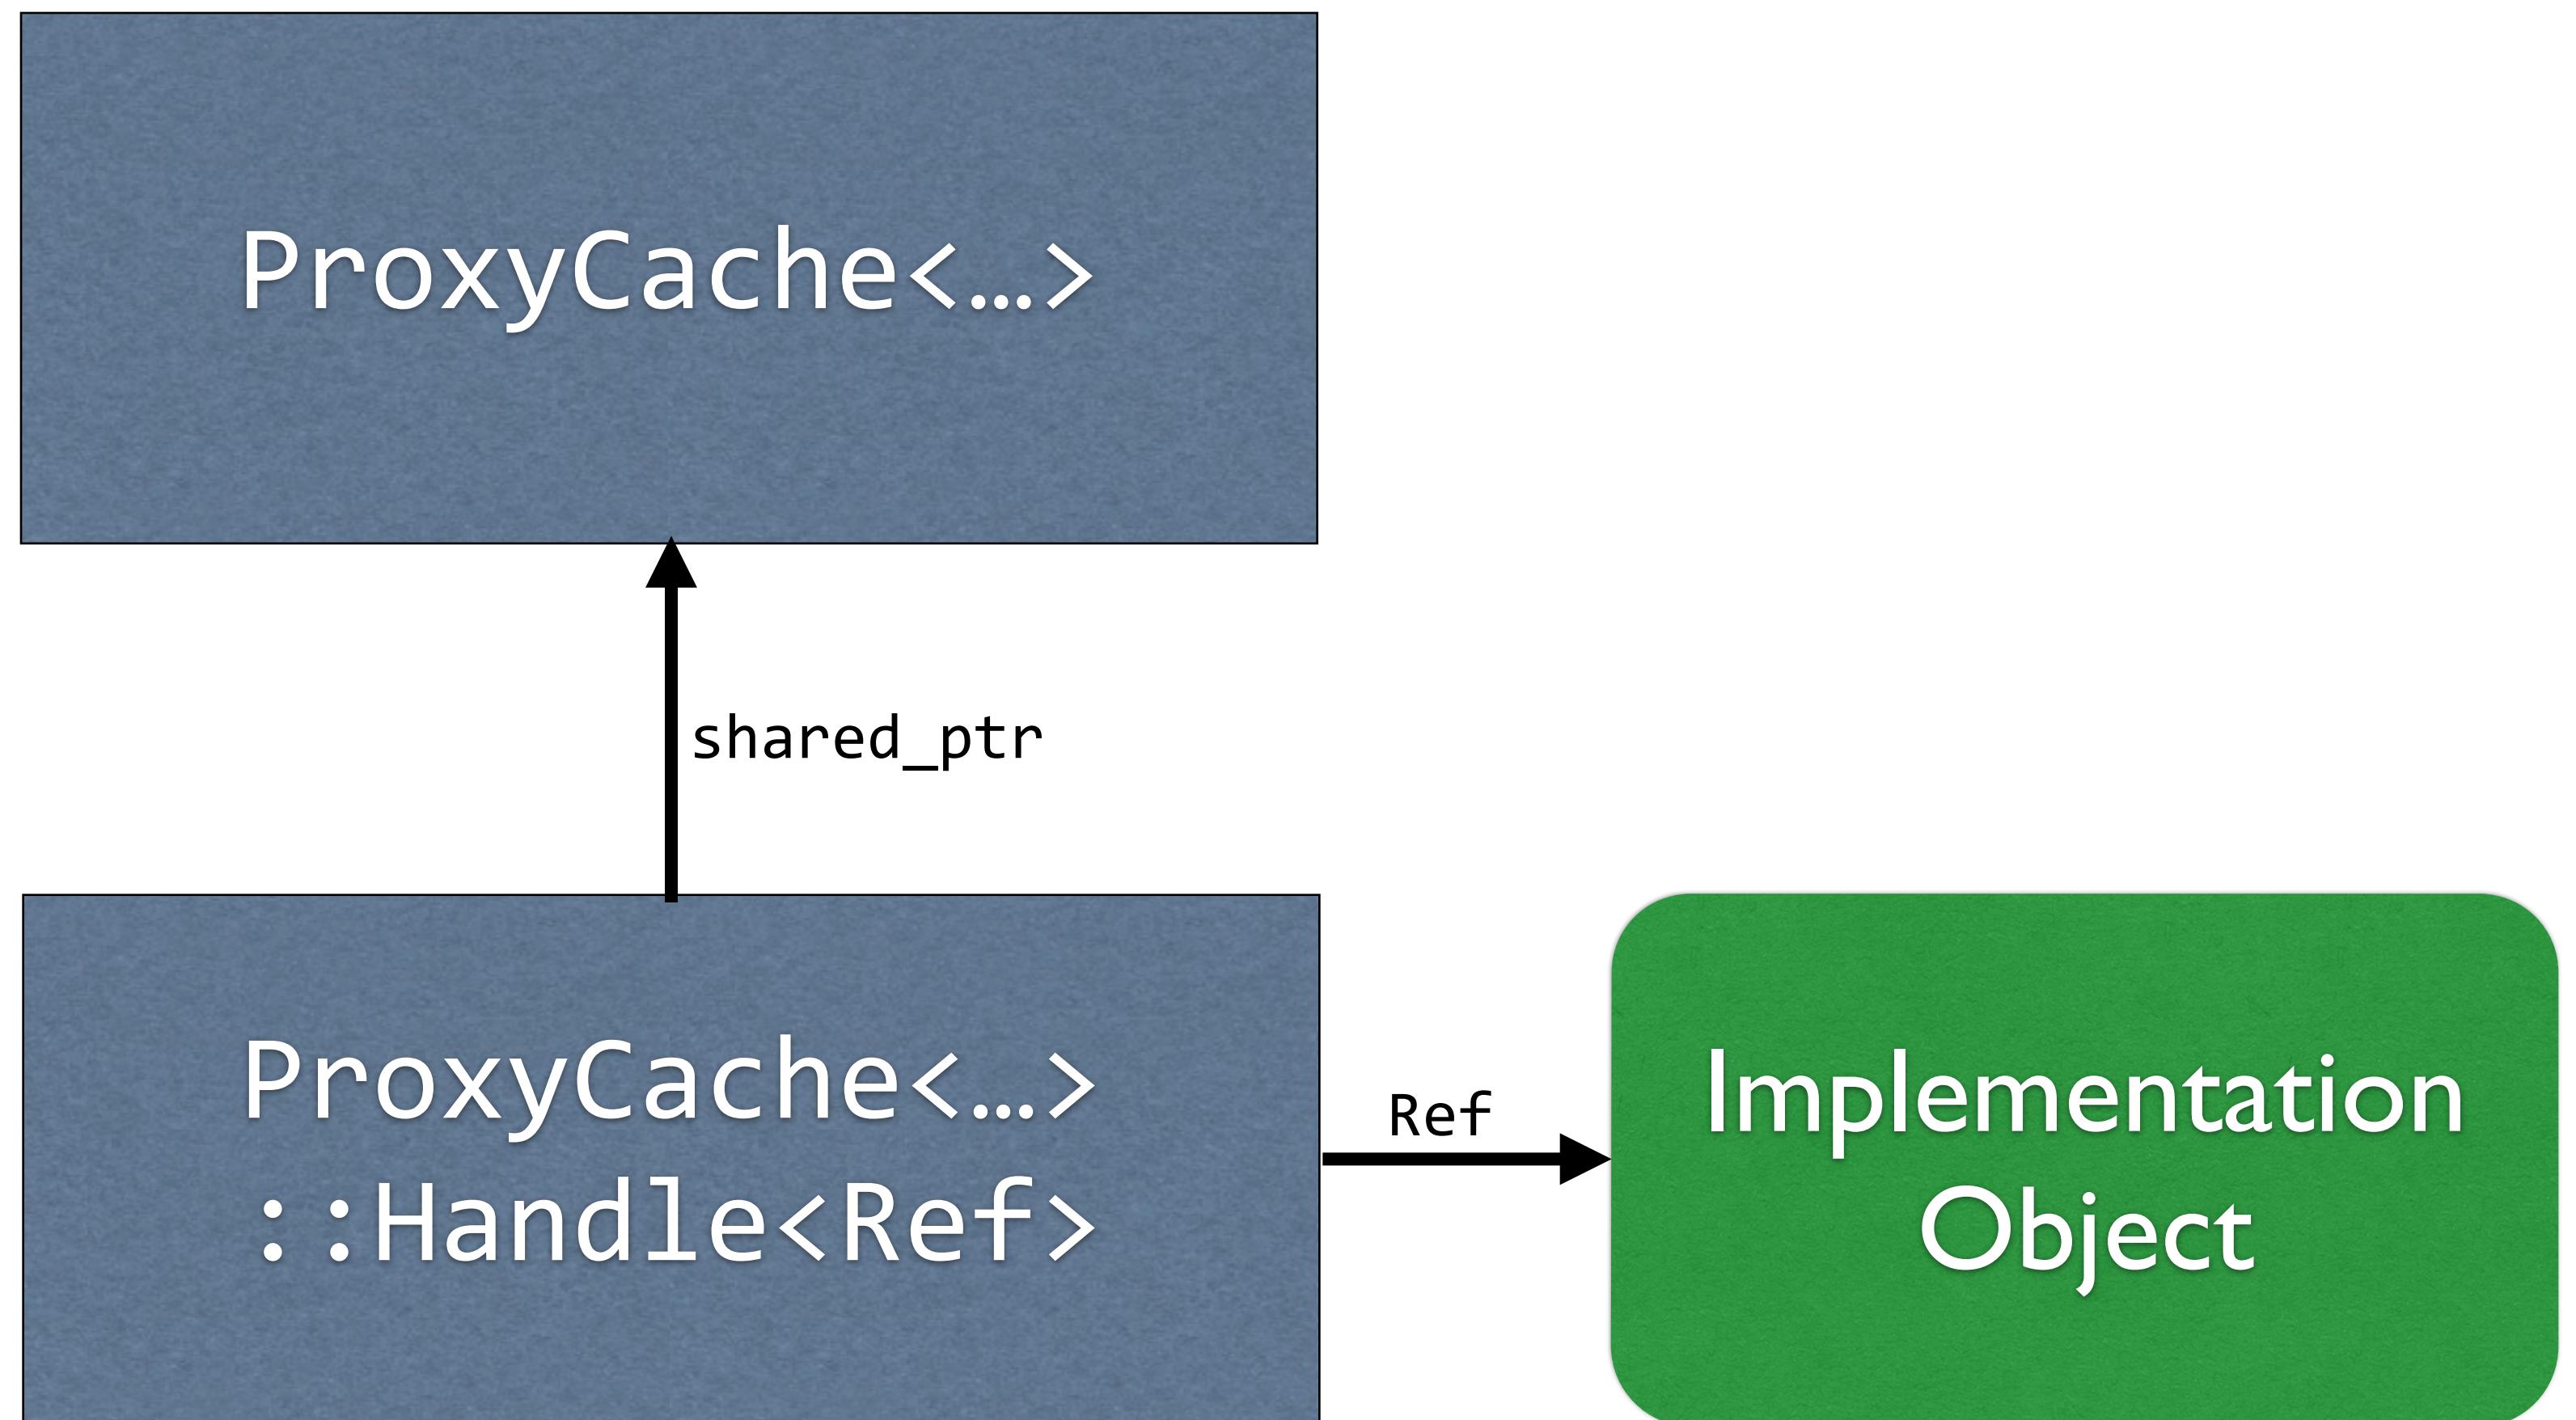
Invariants

- ⌚ Guarantee: no two wrappers for the same object are simultaneously visible
 - ⌚ Not LRU
 - ⌚ Wrappers don't last forever
 - ⌚ *Visible, not in existence*
 - ⌚ Weak-pointer expiry determines when we can make a new wrapper





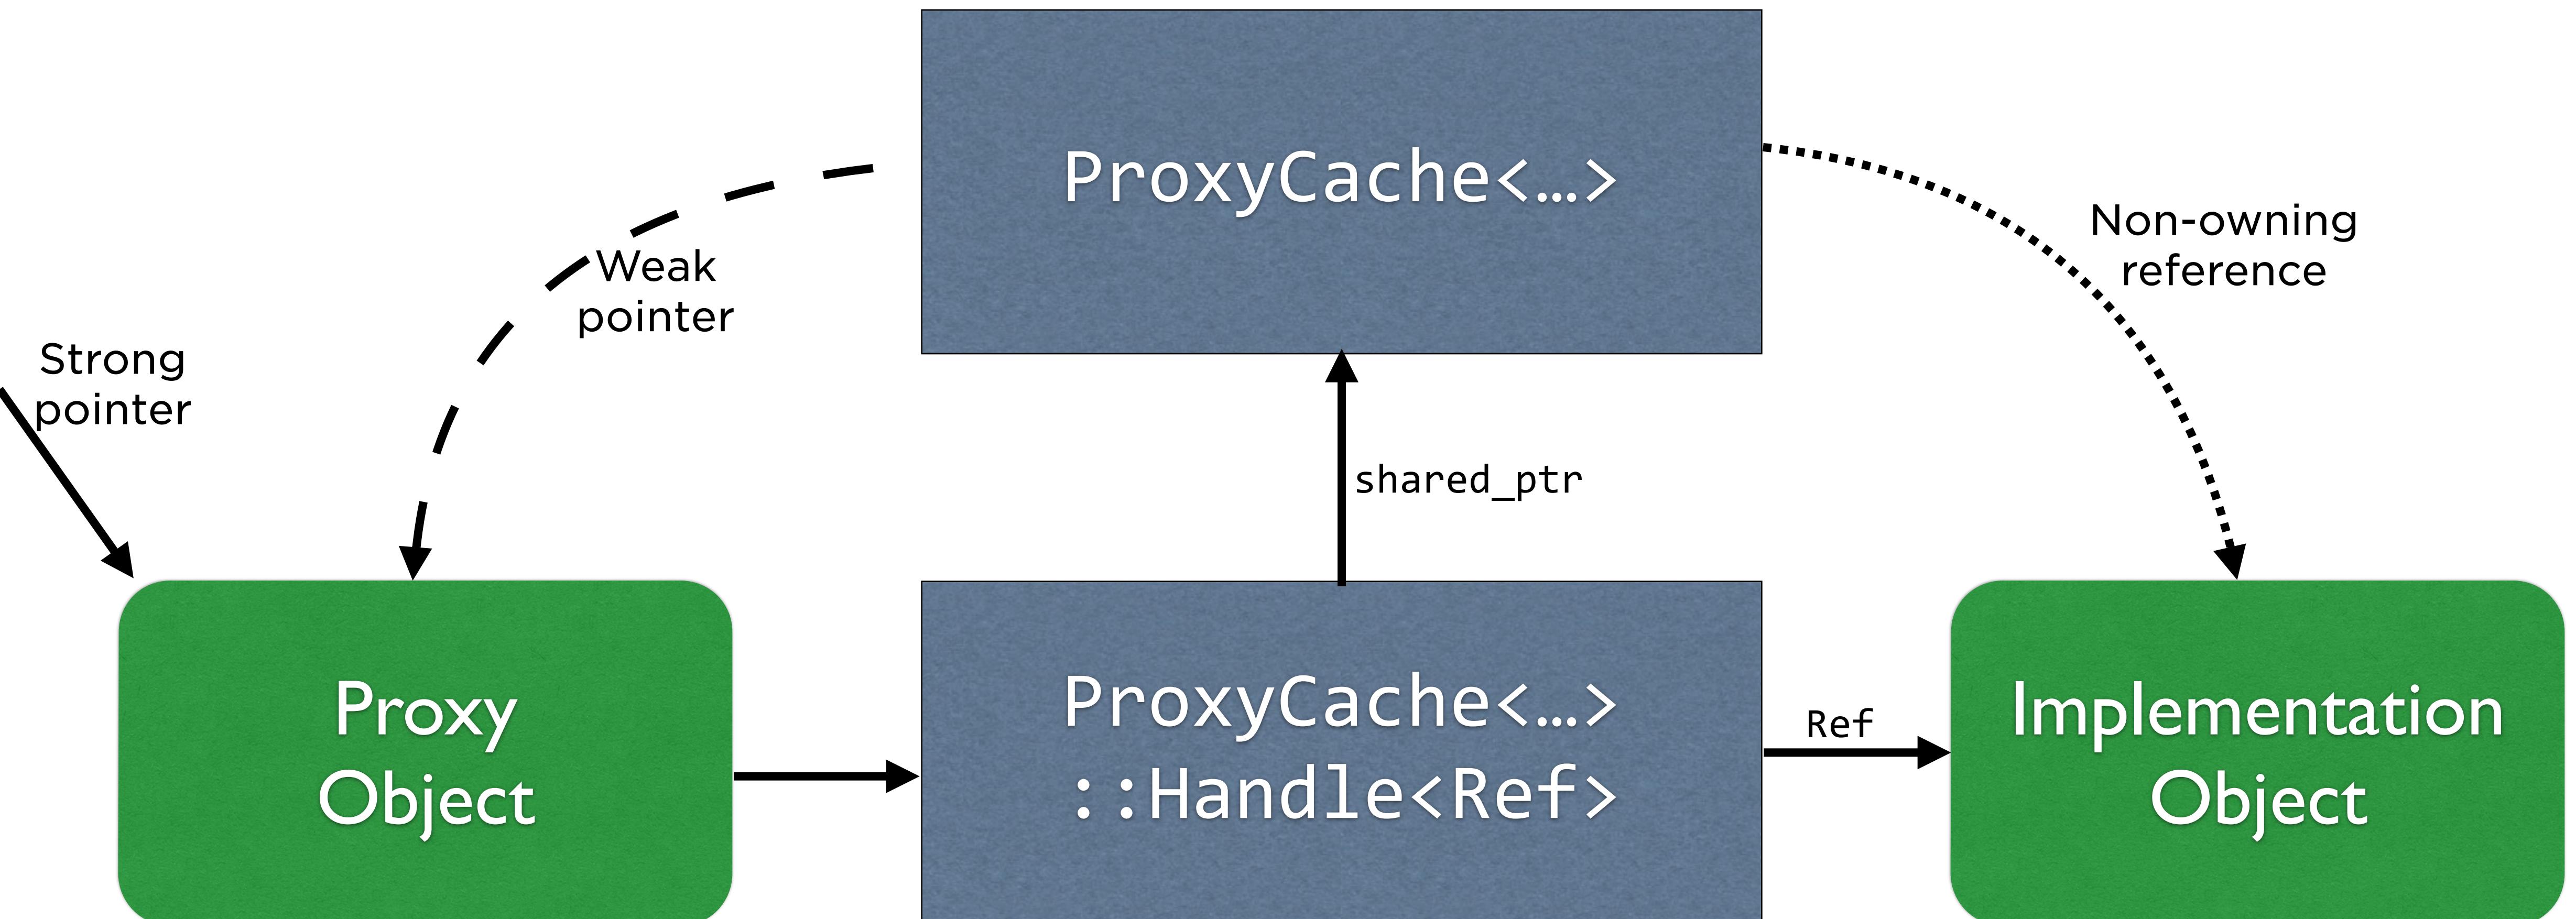


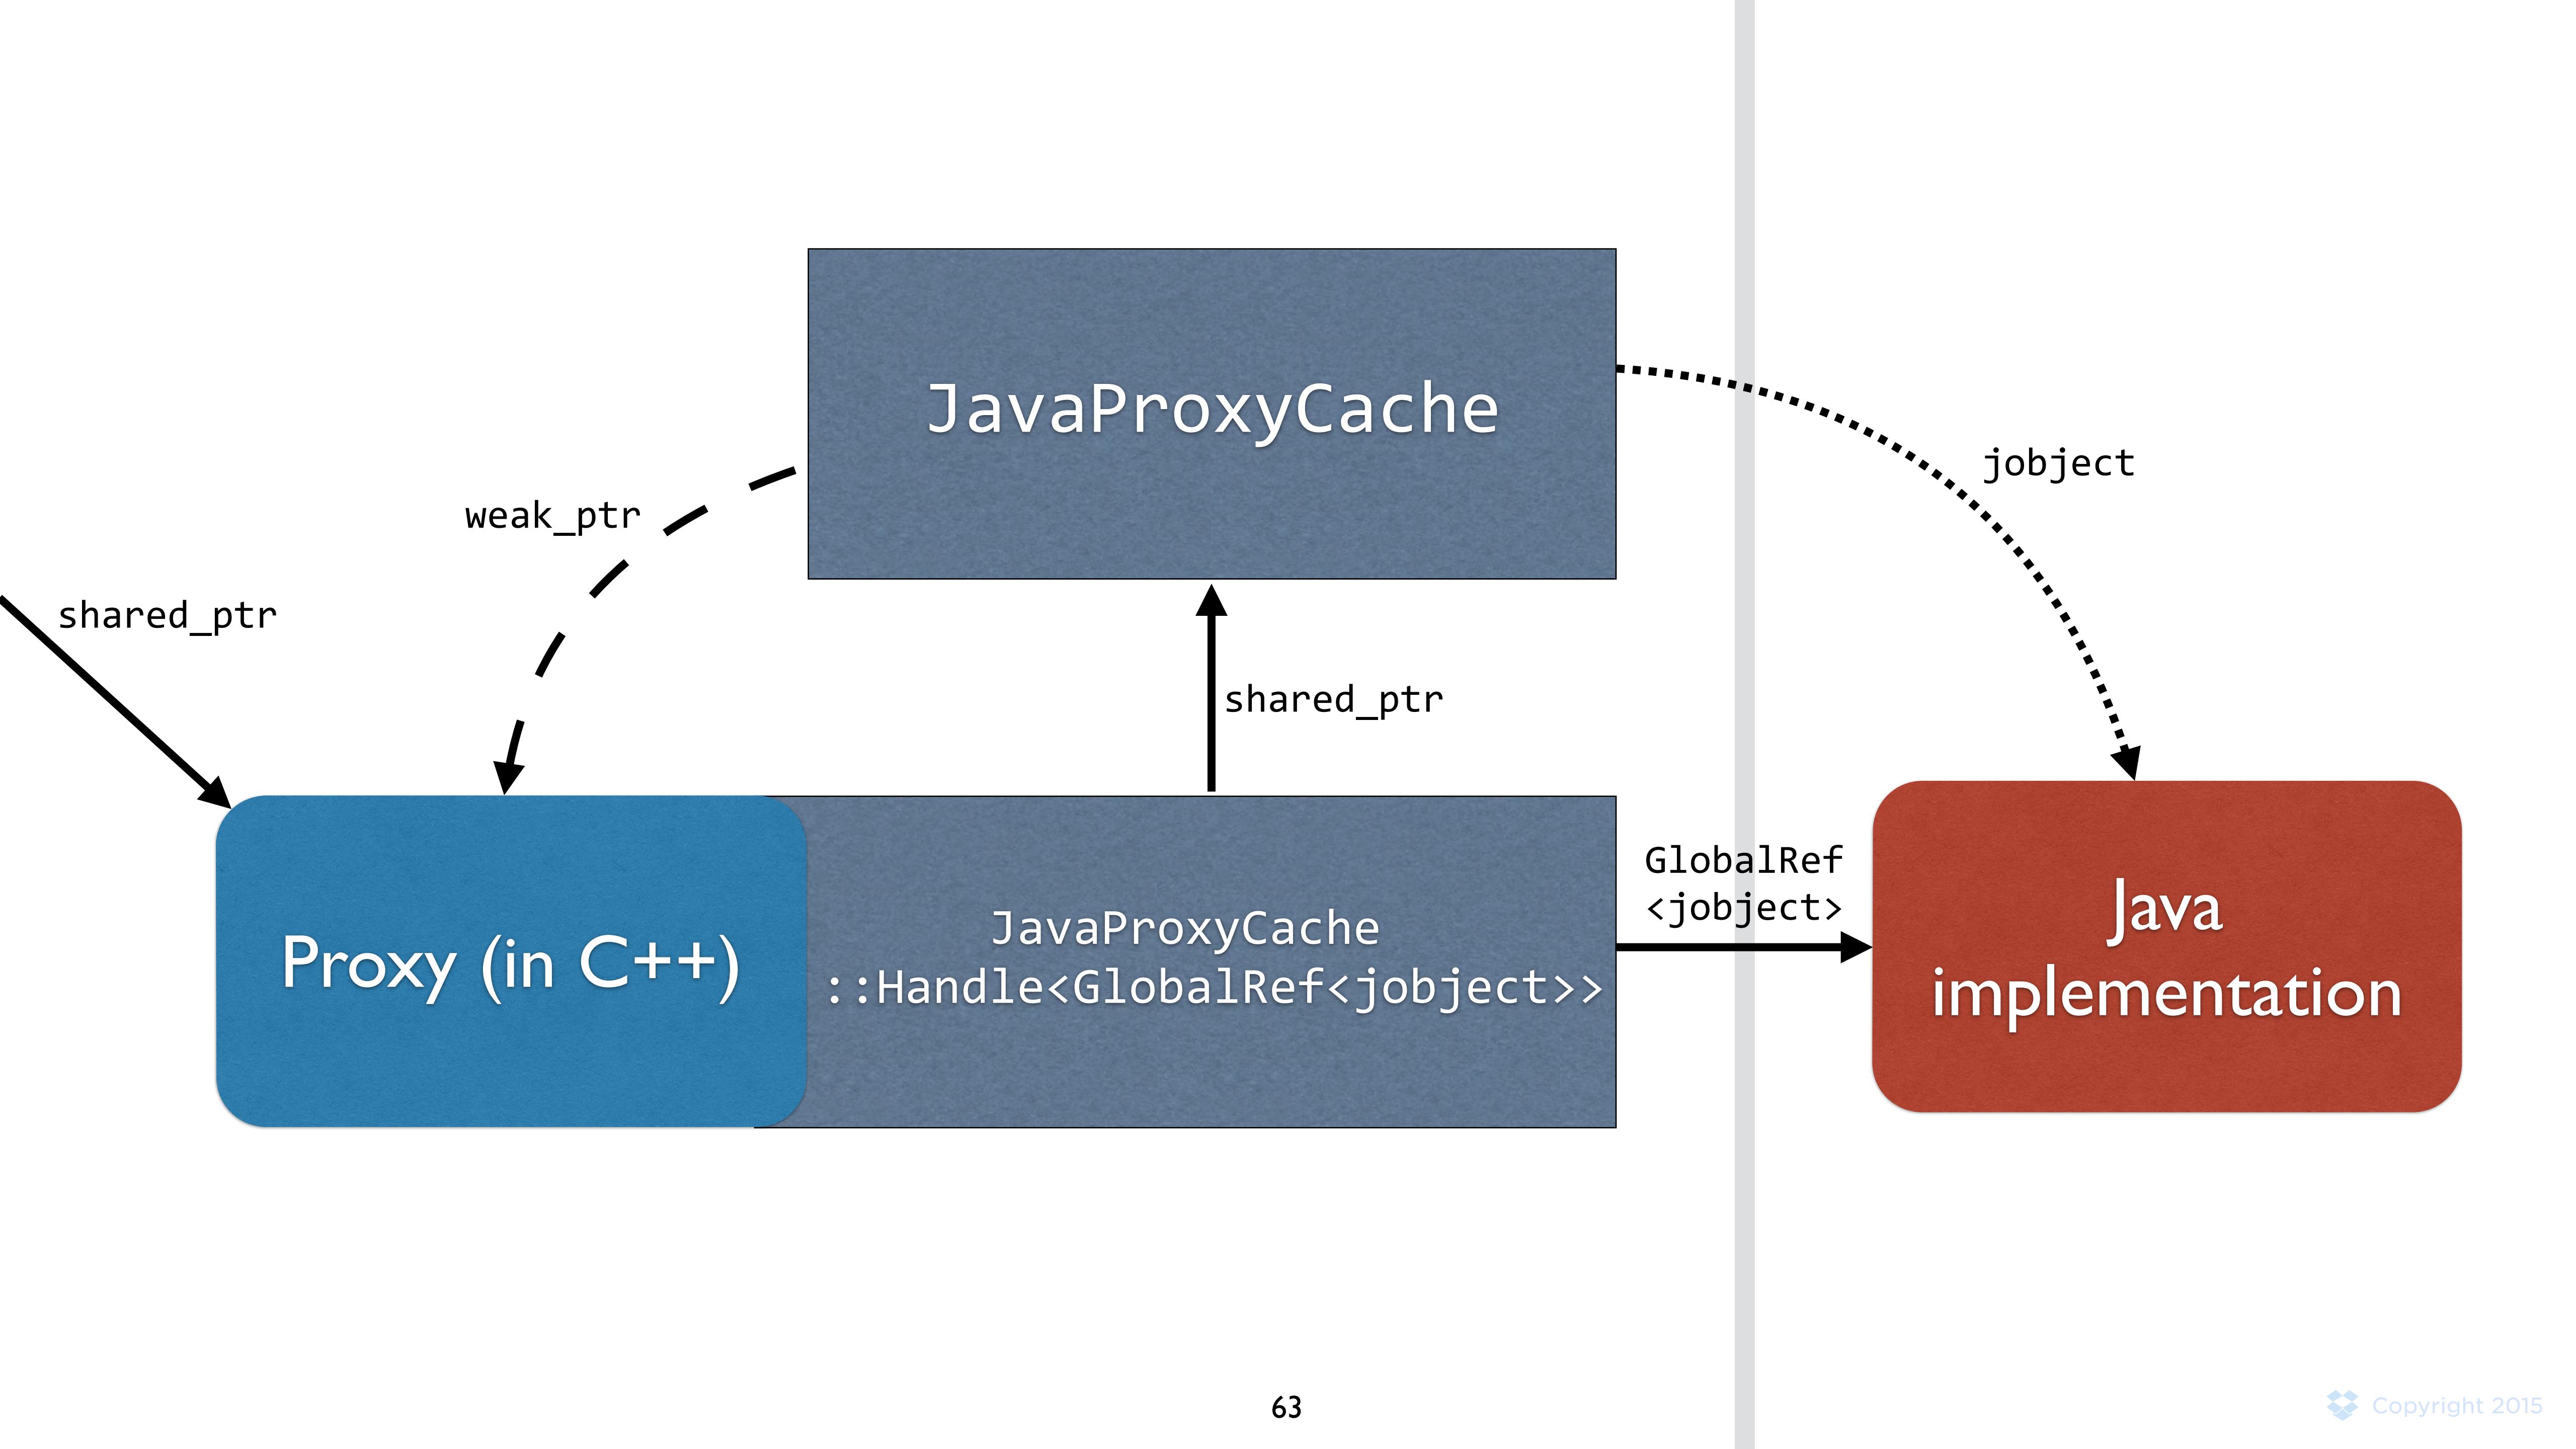


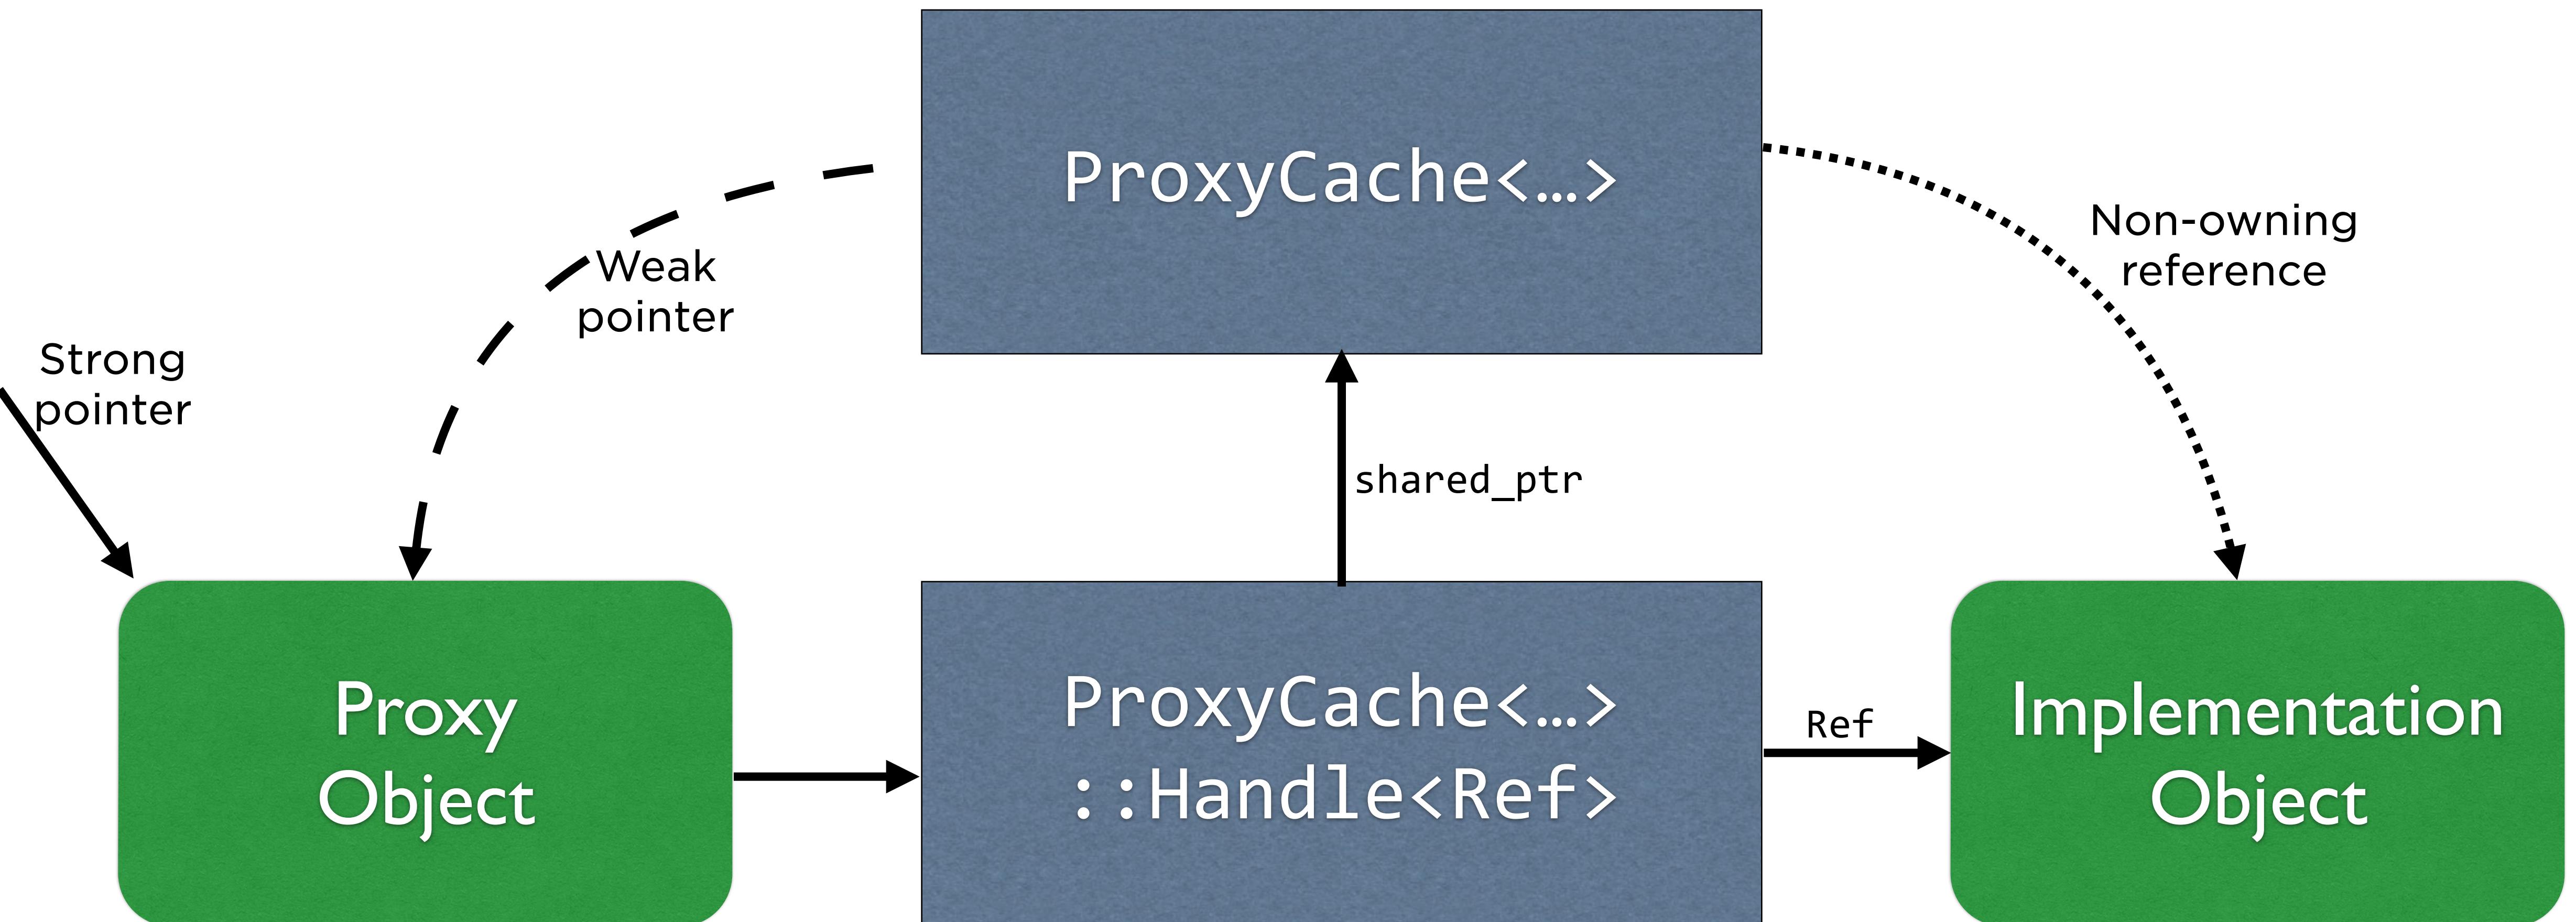
ProxyCache<...>

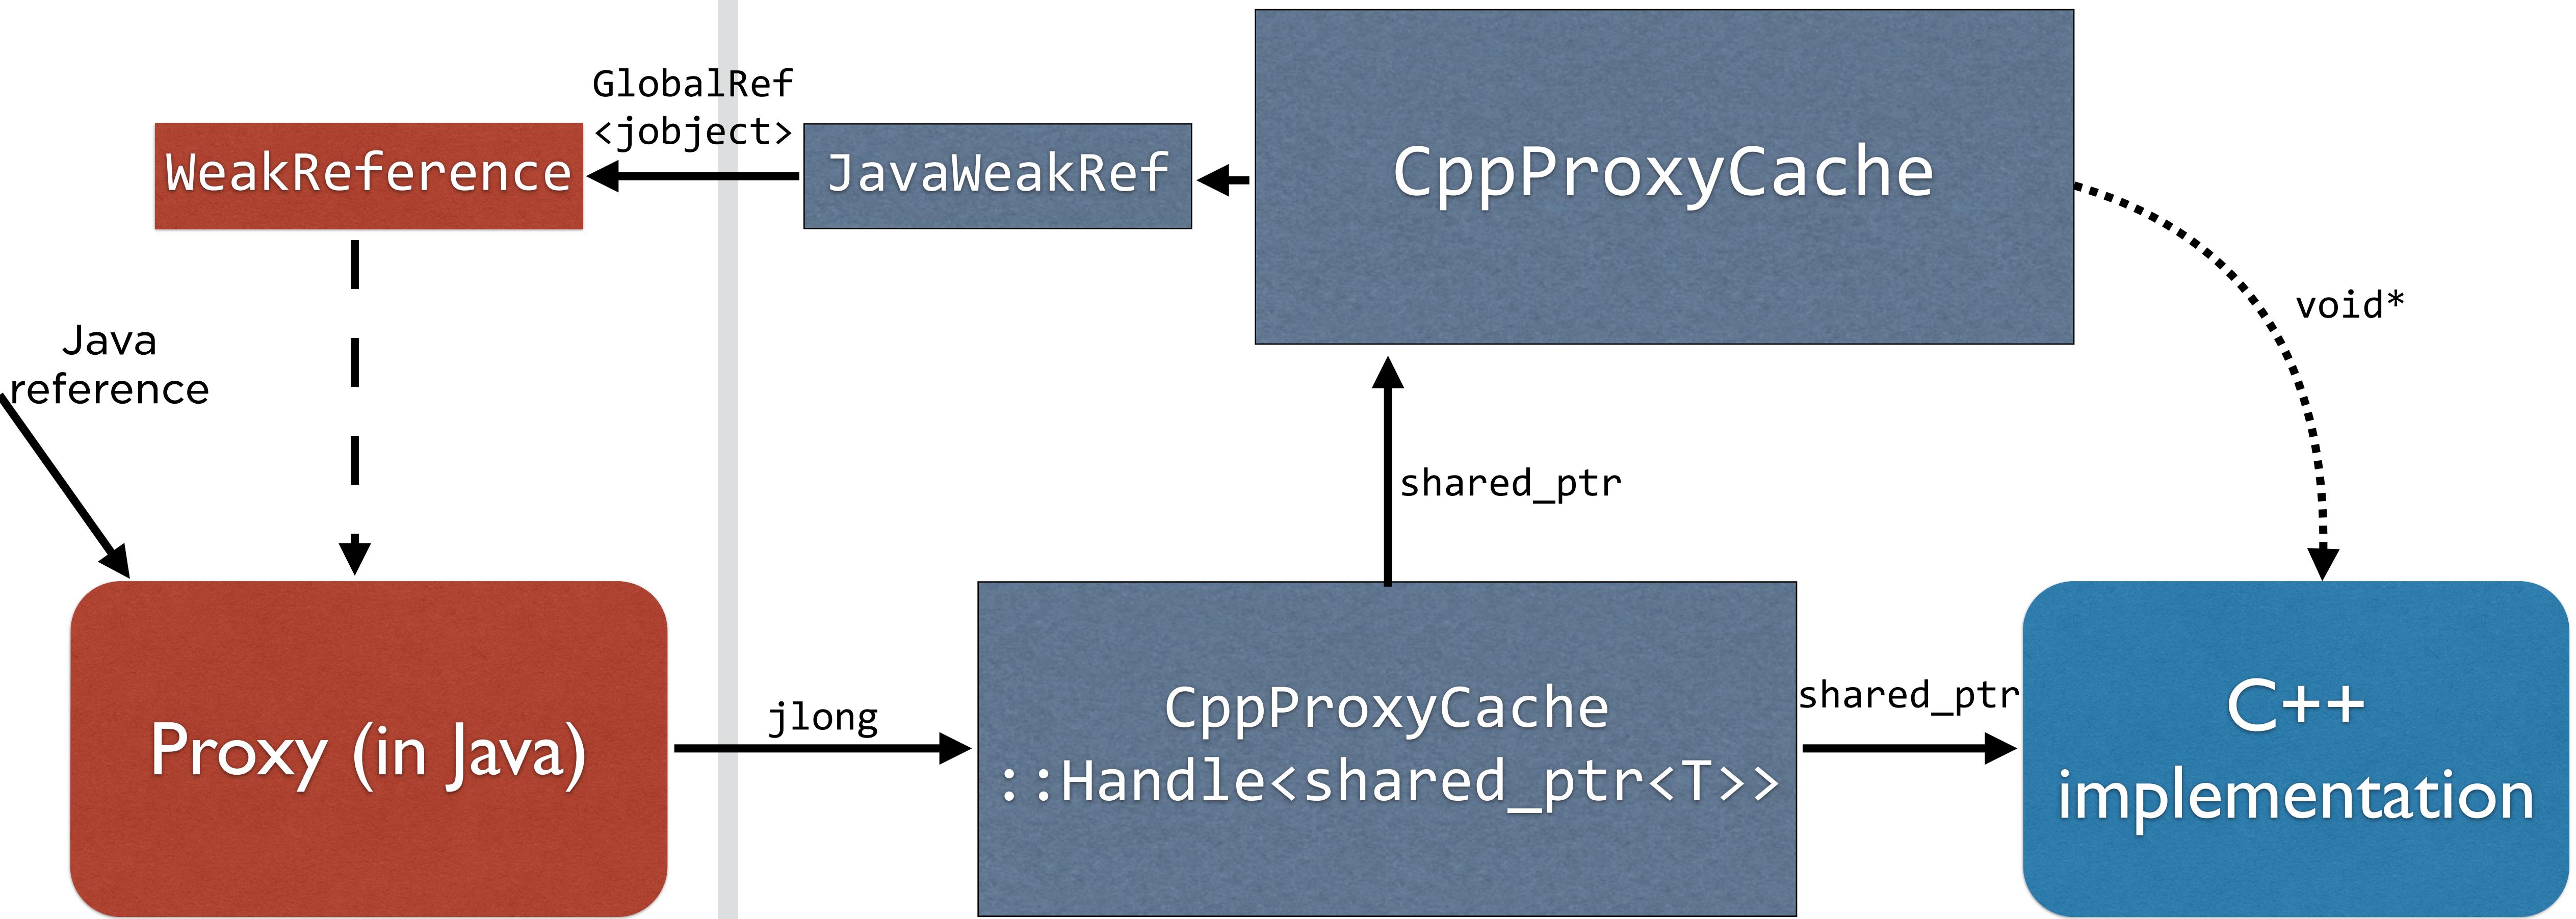
Implementation
Object

ProxyCache<...>









Gotchas

- ⌚ ProxyCache is a global singleton
- ⌚ *But it has to be held by shared_ptr! Why?*
- ⌚ How many kinds of weak reference does Java have?

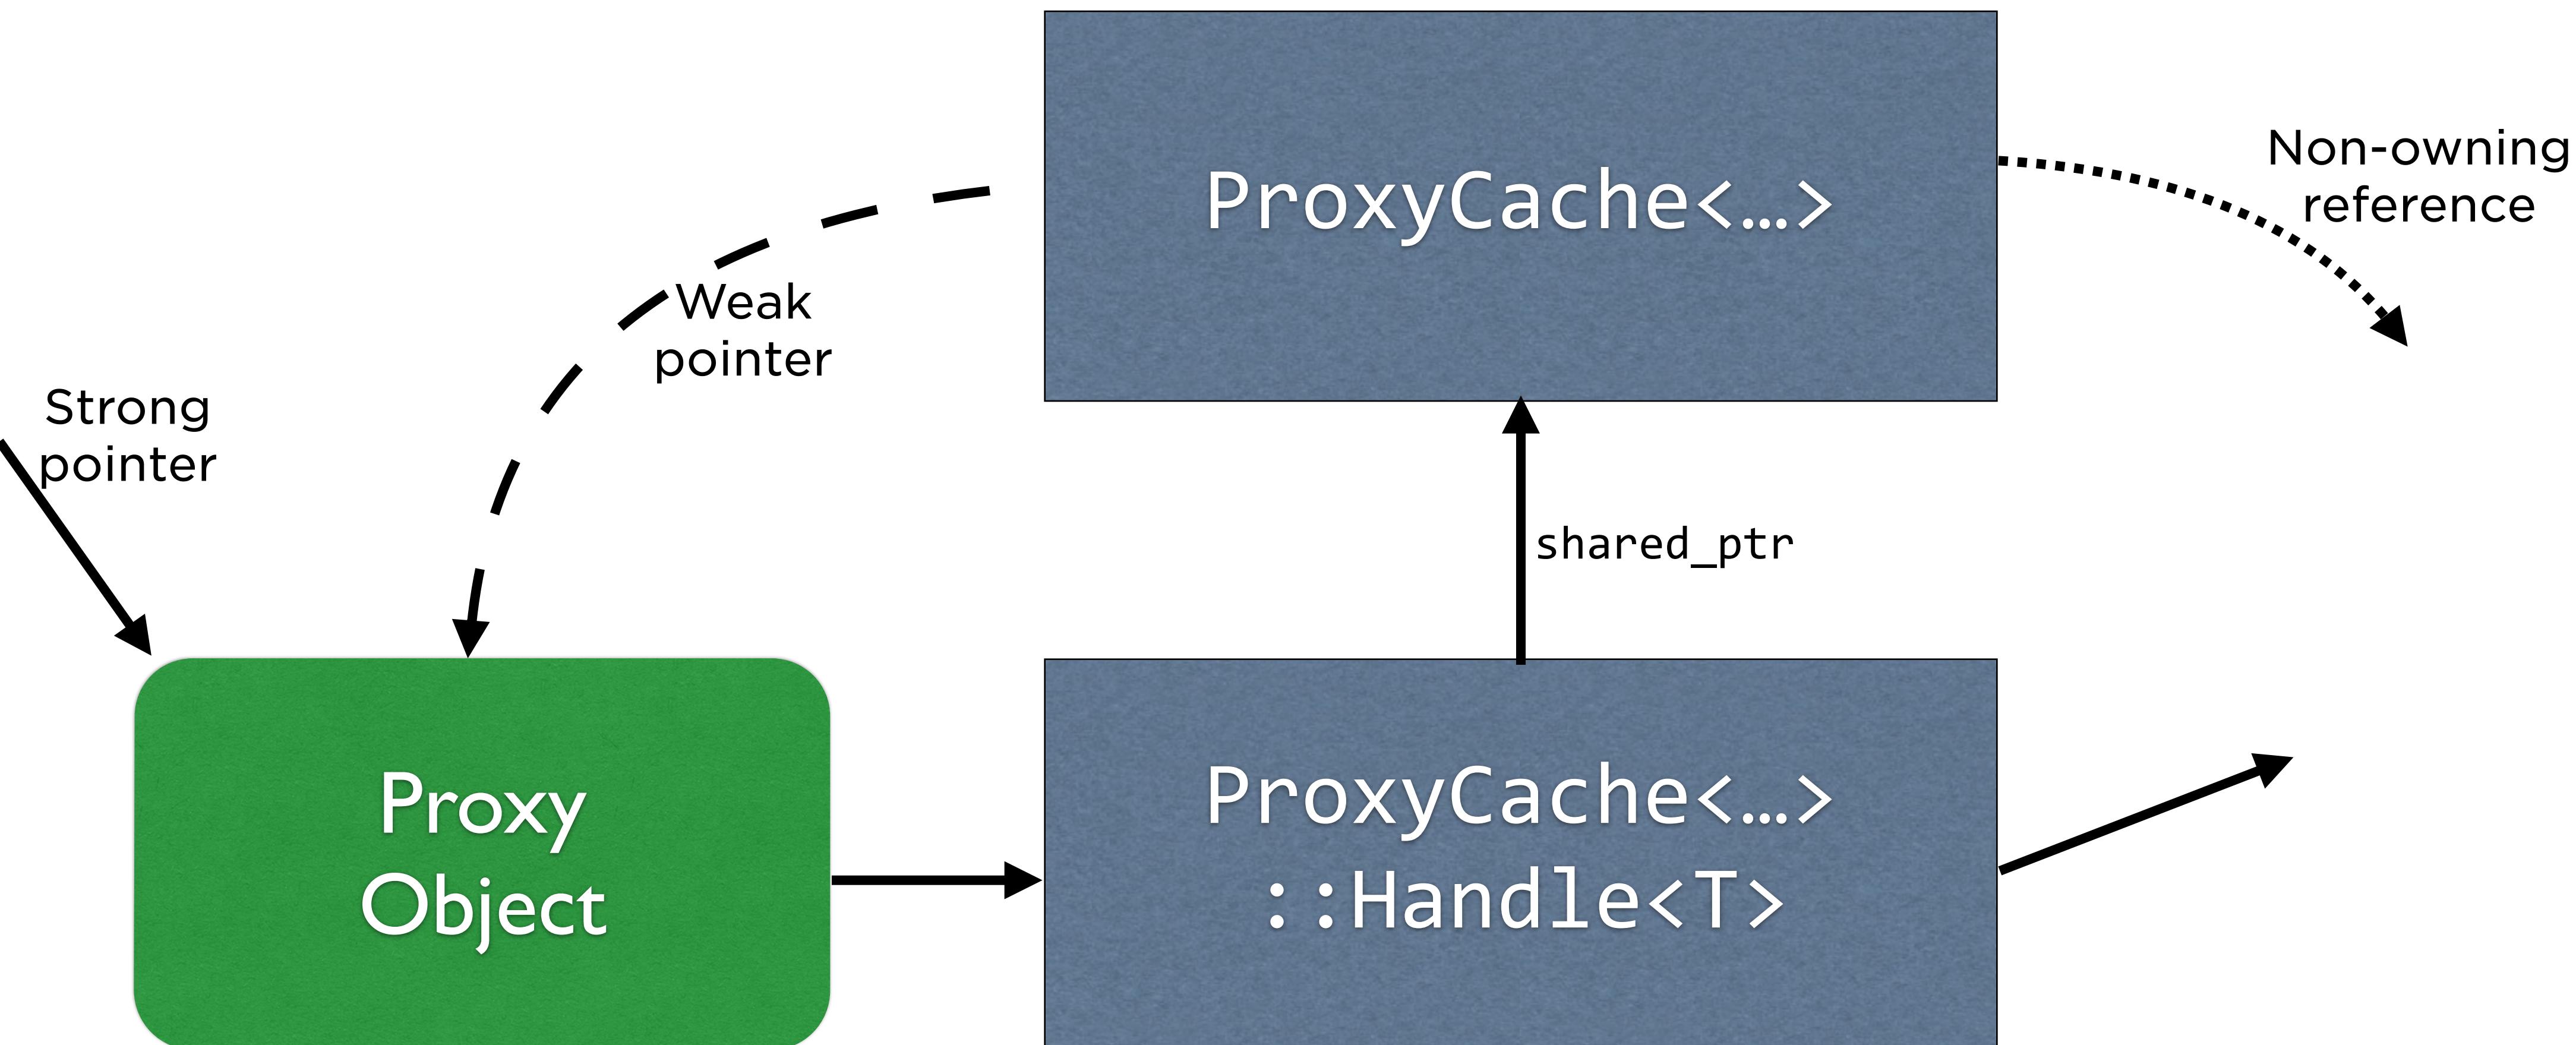
	Java -> C++	C++ -> Java	ObjC -> C++	C++ -> ObjC
UnowningImplPtr	<code>jobject</code>	<code>void *</code>	<code>__unsafe_unretained id</code>	<code>void *</code>
OwningImplPtr	<code>jobject</code>	<code>shared_ptr<void></code>	<code>__strong id</code>	<code>shared_ptr<void></code>
OwningProxyPtr	<code>shared_ptr<void></code>	<code>jobject</code>	<code>shared_ptr<void></code>	<code>__strong id</code>
WeakProxyPtr	<code>weak_ptr<void></code>	<code>JavaWeakRef</code>	<code>weak_ptr<void></code>	<code>__weak id</code>
UnowningImplPtrHash	<code>JavaIdentityHash</code>	<code>std::hash</code>	<code>UnretainedIdHash</code>	<code>std::hash</code>
UnowningImplPtrEqual	<code>JavaIdentityEquals</code>	<code>std::equal_to</code>	<code>std::equal_to</code>	<code>std::equal_to</code>

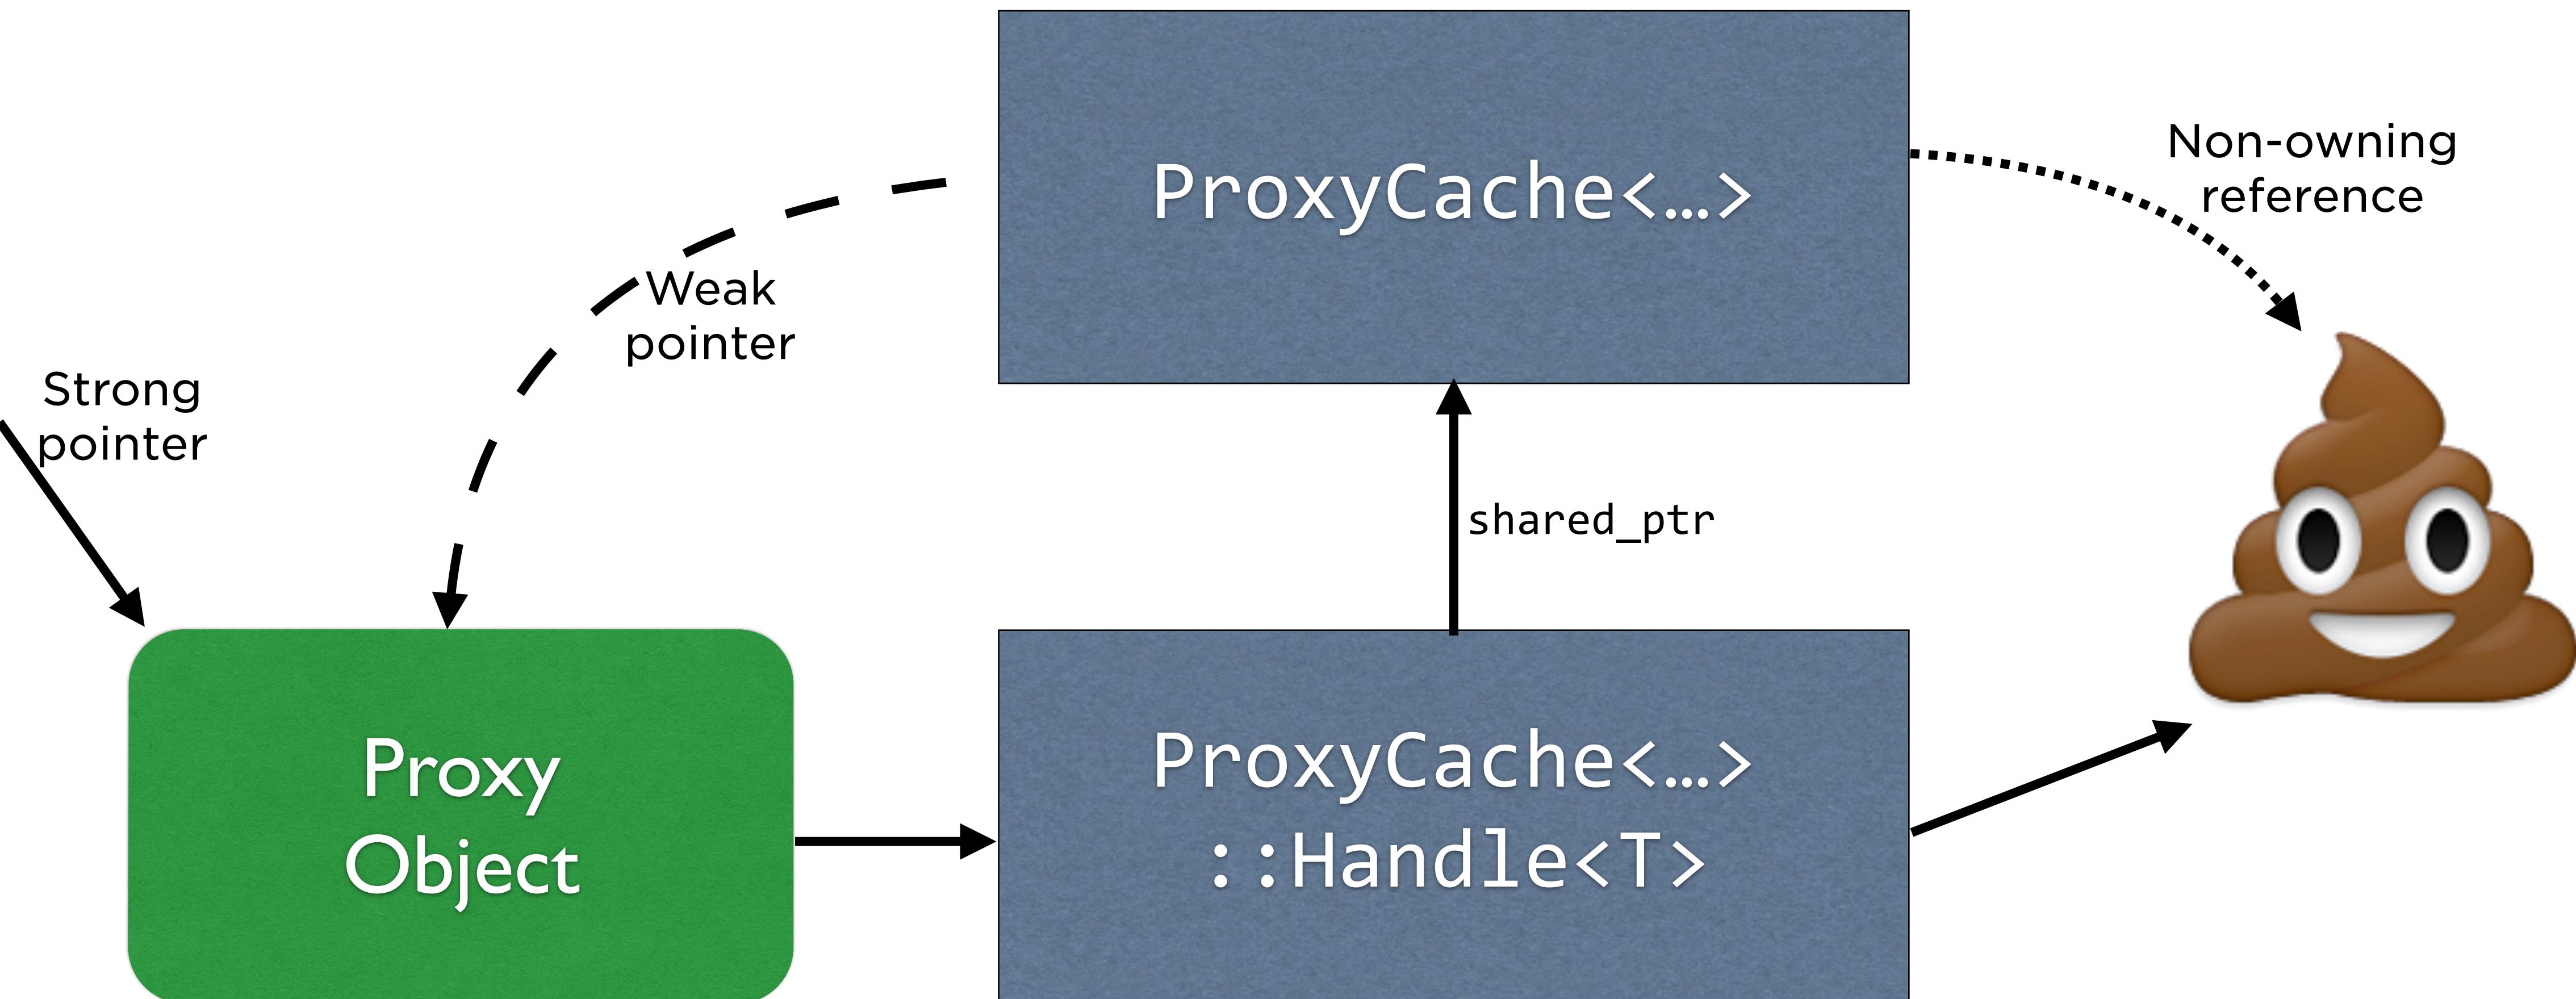
Sync Progress

- ✓ Djinni overview
- ✓ Adding a new language (Python)
- ⌚ Implementation techniques
 - ✓ Proxy caching for identity semantics
 - ⌚ **Nullability**
 - ⌚ Customized Java exception translation

Nullability

- ⌚ A year ago, we didn't handle null at all





Nullability

- ⌚ A year ago, we didn't handle null at all
- ⌚ Changed so null was allowed on interfaces
- ⌚ `optional<interface_type>` becomes `optional<shared_ptr<...>> ?!`
- ⌚ ObjC and Java both have nullability annotations
 - ⌚ `@Nonnull`, `@CheckForNull` (et al), `_Nonnull`, `_Nullable`
 - ⌚ We generated these for data objects and interfaces

Intermediate state

- ➊ Foo: @Nonnull Foo, _Nonnull Foo*, shared_ptr<Foo>
- ➋ optional<Foo>: @CheckForNull Foo, _Nullable Foo*, shared_ptr<Foo>
- ➌ Generated code includes explicit null checks: ObjC and Java annotations are advisory only

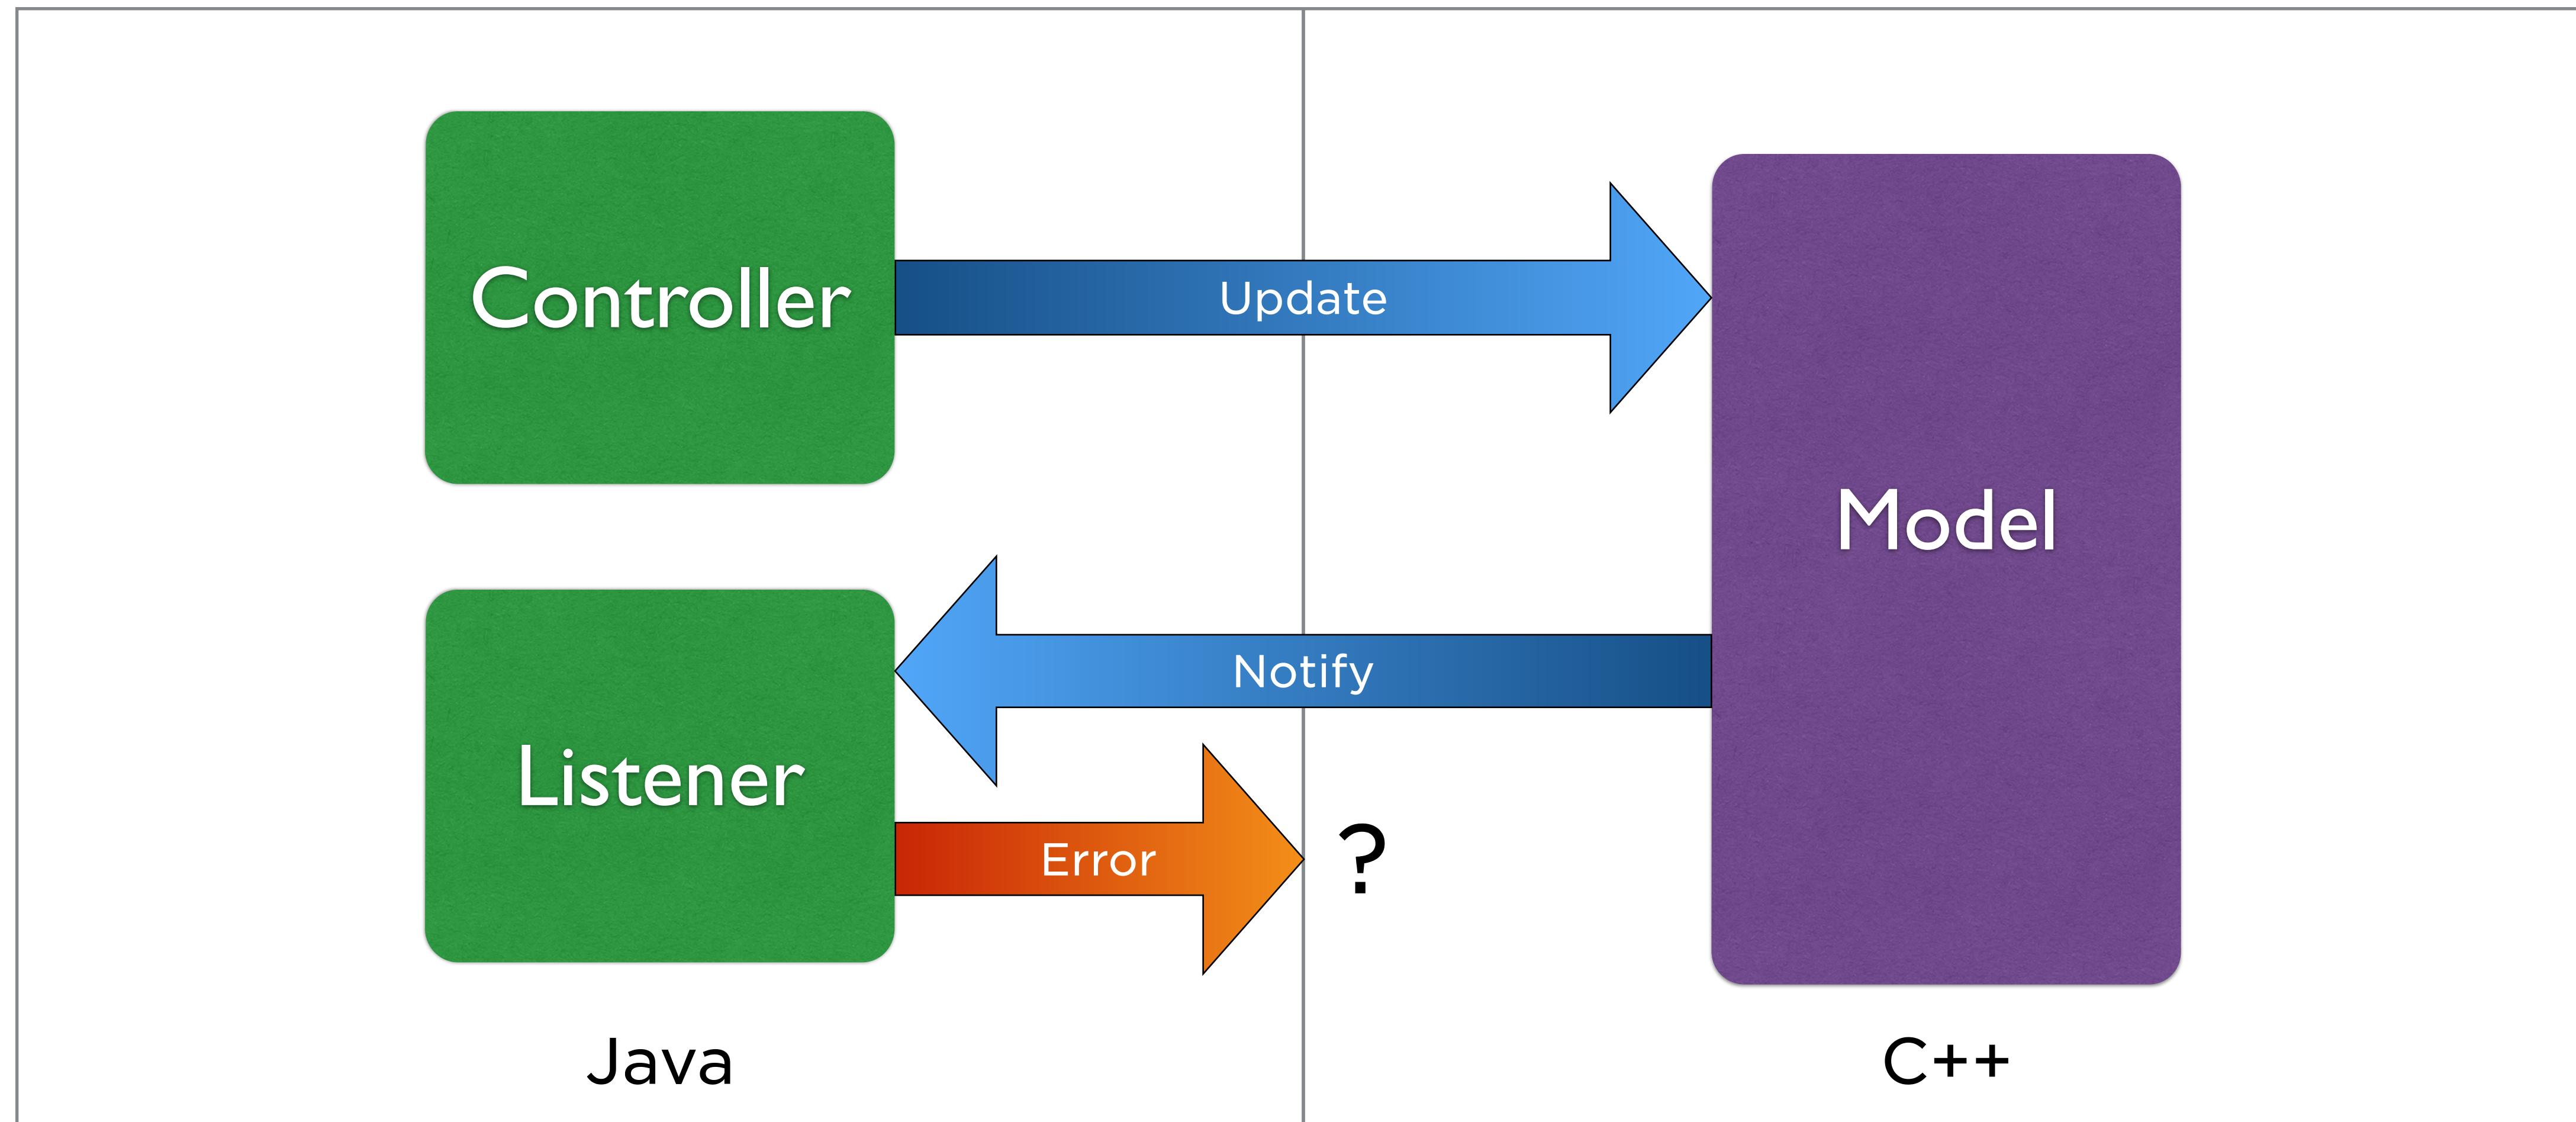
Handling null today

- ➊ Foo is non-nullable, optional<Foo> is nullable
- ➋ Can hook into nn (generating nn_shared_ptr<T>)
- ➌ nn_shared_ptr<T>
 - ➍ A pointer that has *already* been *explicitly* checked for null
 - ➎ (GSL's not_null inserts implicit checks)
- ➏ For more, see my talk tomorrow, 2:30pm

Sync Progress

- ✓ Djinni overview
- ✓ Adding a new language (Python)
- ⌚ Implementation techniques
 - ✓ Proxy caching for identity semantics
 - ✓ Nullability
- ⌚ **Customized Java exception translation**

Exception Translation



JNI Exception Model

- ➊ Each thread can have a “pending” exception
- ➋ Java side is done for us
 - ➌ Catch and Set before return from Java
 - ➌ Check and throw after return to Java
- ➌ Our generated code must do the other side
 - ➌ Check and throw before after return from Java
 - ➌ Catch and set before return to Java

Djinni's Default Translation

- ➊ After return from Java (Java → C++)
 - ➊ Get and clear the pending Java exception
 - ➋ Throw `djinni::jni_exception` holding it
- ➋ Before return to Java (C++ → Java)
 - ➊ Catch `jni_exception`
 - ➋ Extract the Java exception
 - ➋ Set it as pending again and return
 - ➌ For any other C++ exception, create a `RuntimeException`

What if that's not good enough?

- ⌚ This is free, but simplistic
- ⌚ What if you have your own custom exception types?
- ⌚ What if you want stack traces from C++?

Pluggable Functions

- ➊ You can replace exception translation with your own version
- ➋ Helper functions defined with `__attribute__((weak))`
- ➌ Linker will prefer your definition, if you provide it

- ➍ For example:
 - ➎ You have a `my_security_exception` you use in C++
 - ➏ Should translate to/from `SecurityException` in Java

Translating Java → C++

```
void jniThrowCppFromJavaException(JNIEnv * env,  
                                  jthrowable java_exception);
```

- ⌚ Called when a Java exception makes it to C++
- ⌚ Should always throw some C++ exception

```
void jniThrowCppFromJavaException(JNIEnv * env, jthrowable java_exception) {
    const SecurityExceptionClassInfo & ci =
        djinni::JniClass<SecurityExceptionClassInfo>::get();

    if (env->IsInstanceOf(java_exception, ci.clazz.get())) {
        LocalRef<jstring> jmessage {
            env, (jstring)env->CallObjectMethod(java_exception, ci.getMessage) };

        throw my_security_exception(jniUTF8FromString(env, jmessage.get()));
    }

    throw jni_exception { env, java_exception };
}
```

```
struct SecurityExceptionClassInfo {
    const GlobalRef<jclass> clazz =
        djinni::jniFindClass("java/lang/SecurityException");

    const jmethodID ctor =
        djinni::jniGetMethodID(clazz.get(),
                               "<init>",
                               "(Ljava/lang/String;)V");

    const jmethodID getMessage =
        djinni::jniGetMethodID(clazz.get(),
                               "getMessage",
                               "()Ljava/lang/String;");
};
```

Translating C++ → Java

```
void jniSetPendingFromCurrent(JNIEnv * env, const char * ctx);
```

- ⌚ Called when a C++ exception makes it to Java
- ⌚ Called inside a catch block
- ⌚ Should always set the pending exception and return

```
void jniSetPendingFromCurrent(JNIEnv * env, const char * ctx) noexcept {
    try {
        throw;
    } catch (const my_security_exception & e) {
        const SecurityExceptionClassInfo & ci =
            djinni::JniClass<SecurityExceptionClassInfo>::get();

        LocalRef<jstring> jmessage { env, jniStringFromUTF8(env, e.what()) };

        LocalRef<jthrowable> jexception { env,
            (jthrowable)env->NewObject(ci.clazz.get(), ci.ctor, jmessage.get()) };

        env->Throw(jexception);
        return;
    } catch (const std::exception & e) {
        jniDefaultSetPendingFromCurrent(env, ctx);
    }
}
```

Add Stack Traces

- ⌚ Can be built using the same two helpers
- ⌚ Sample code is too big for slides
- ⌚ I'll give the key points
- ⌚ I assume Java is always the root (last to catch)
 - ⌚ It needs to see the full stack trace

Enabling Factors

- ⌚ Assume you have a way to generate stack traces in C++
- ⌚ Two things we need to leverage here from Java Throwable
 - ⌚ They can have a cause (another Throwable)
 - ⌚ You can override the stack trace (StackTraceElement [])

Define some helpers

- ➊ A C++ a new exception type containing:
 - ➊ C++ stack trace
 - ➋ `jni_exception (cause)`
- ➋ A Java “builder” type for building exceptions
 - ➊ Set message
 - ➋ Set cause
 - ➌ Add stack trace elements (1 at a time)

Java → C++

- ④ Store original Java exception inside your C++ exception
- ④ Fill in the C++ backtrace
- ④ Throw and let it propagate as normal

C++ → Java

- ➊ Use your builder to create a cause chain
 - ➊ New Java exception (with default stack trace)
 - ➋ ... caused by pseudo-C++ exception (custom stack trace)
 - ➋ ... caused by original Java exception
- ➋ Set the newly-built exception as pending and return

Other new stuff this year

- ➊ New types: f32, date
- ➋ External types
- ➌ Example extensions
 - ➍ Thread creation service
 - ➎ JNI Exception customization
- ➏ Trackable dependencies
- ➐ JetBrains plug-in for .djinni files

Community Links

- ➊ Djinni: github.com/dropbox/djinni
- ➋ Talk from CppCon 2014:bit.ly/djinnitalk, bit.ly/djinnivideo
- ➌ Today's Talk: bit.ly/djinnitalk2, bit.ly/djinnivideo2 (soon)
- ➍ Slack Community: mobilecpp.herokuapp.com
- ➎ Developer Tutorials: mobilecpptutorials.com
- ➏ Contact us: atwyman@dropbox.com, j4.cbo@dropbox.com
- ➐ Also we're hiring! dropbox.com/jobs

